

Java's Untapped Potential: High-Performance In-Memory Data Processing with Java

Markus Kett

CEO at Cyrock.AI, Inc.

m.kett@cyrock.ai

[LinkedIn: /MarkusKett](https://www.linkedin.com/in/MarkusKett)

[X:@MarkusKett](https://twitter.com/MarkusKett)





Markus Kett

CEO at Cyrock.AI

X: @MarkusKett

LinkedIn: markuskett

Email:

m.kett@microstream.one

m.kett@cyrock.ai

- XDEV Java IDE
- RapidClipse - Visual Java IDE
- JBoss Hibernate Tools Improvements
- JPA-SQL
- Eclipse Serializer
- EclipseStore
- Eclipse Data Grid
- Project Helidon
- Micronaut

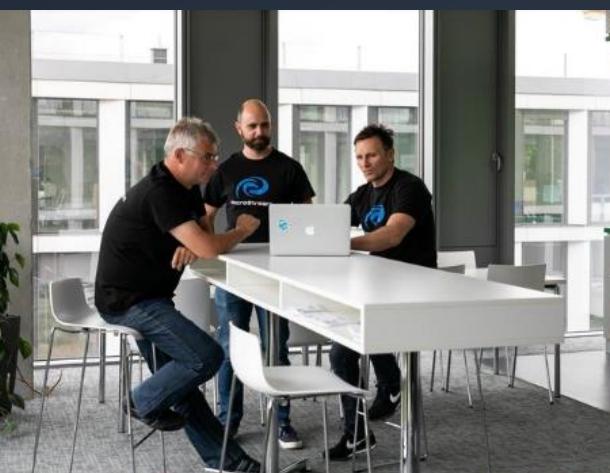
Founder of JAVAPRO Magazine

Editor in Chief

**Founder of JCON Conference Series
Speaker, Author**



The Team Behind Eclipse Data Grid & Cyrock.AI



Eclipse Foundation: More Than 100 Popular Open-Source & Standardization Projects are Developed by Numerous Companies.

We are Managing 3 Eclipse Projects:
Eclipse Data Grid, EclipseStore, and Eclipse Serializer.



projects.eclipse.org



In-Memory Computing is Everywhere



Latency Numbers Every Programmer Should Know



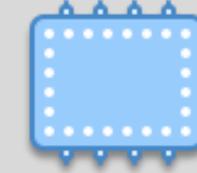
HDD
~ 9 ms
~ 9,000,000 ns



SSD HDD
~ 250 µs
~ 250,000 ns



DRAM
~ 65 ns



CPU Cache
< 1 ns

~ **100,000x faster than HDD**
~ **4,500x faster than SSD**

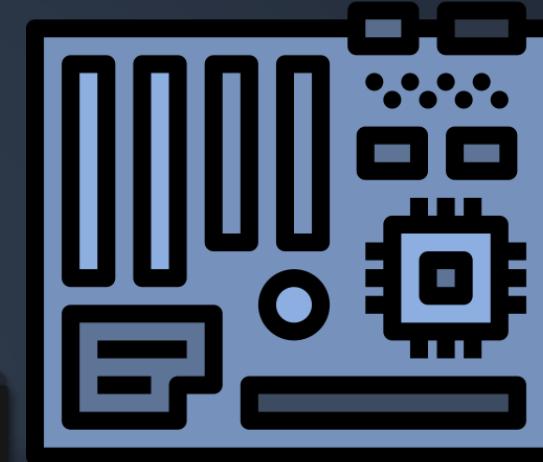
In-Memory Computing: Code Executed Inside a System.



Serverside In-Memory Computing:
Code Executed On the Mainboard.



HDD
~ 9 ms
~ 9,000,000 ns



In-Memory

DRAM
~ 65 ns

~ 100,000x faster than HDD
~ 4,500x faster than SSD



Distributed In-Memory Systems are 1000x Slower



**Network is Factor
1000x - 1,000,000x
Slower Than DRAM.**



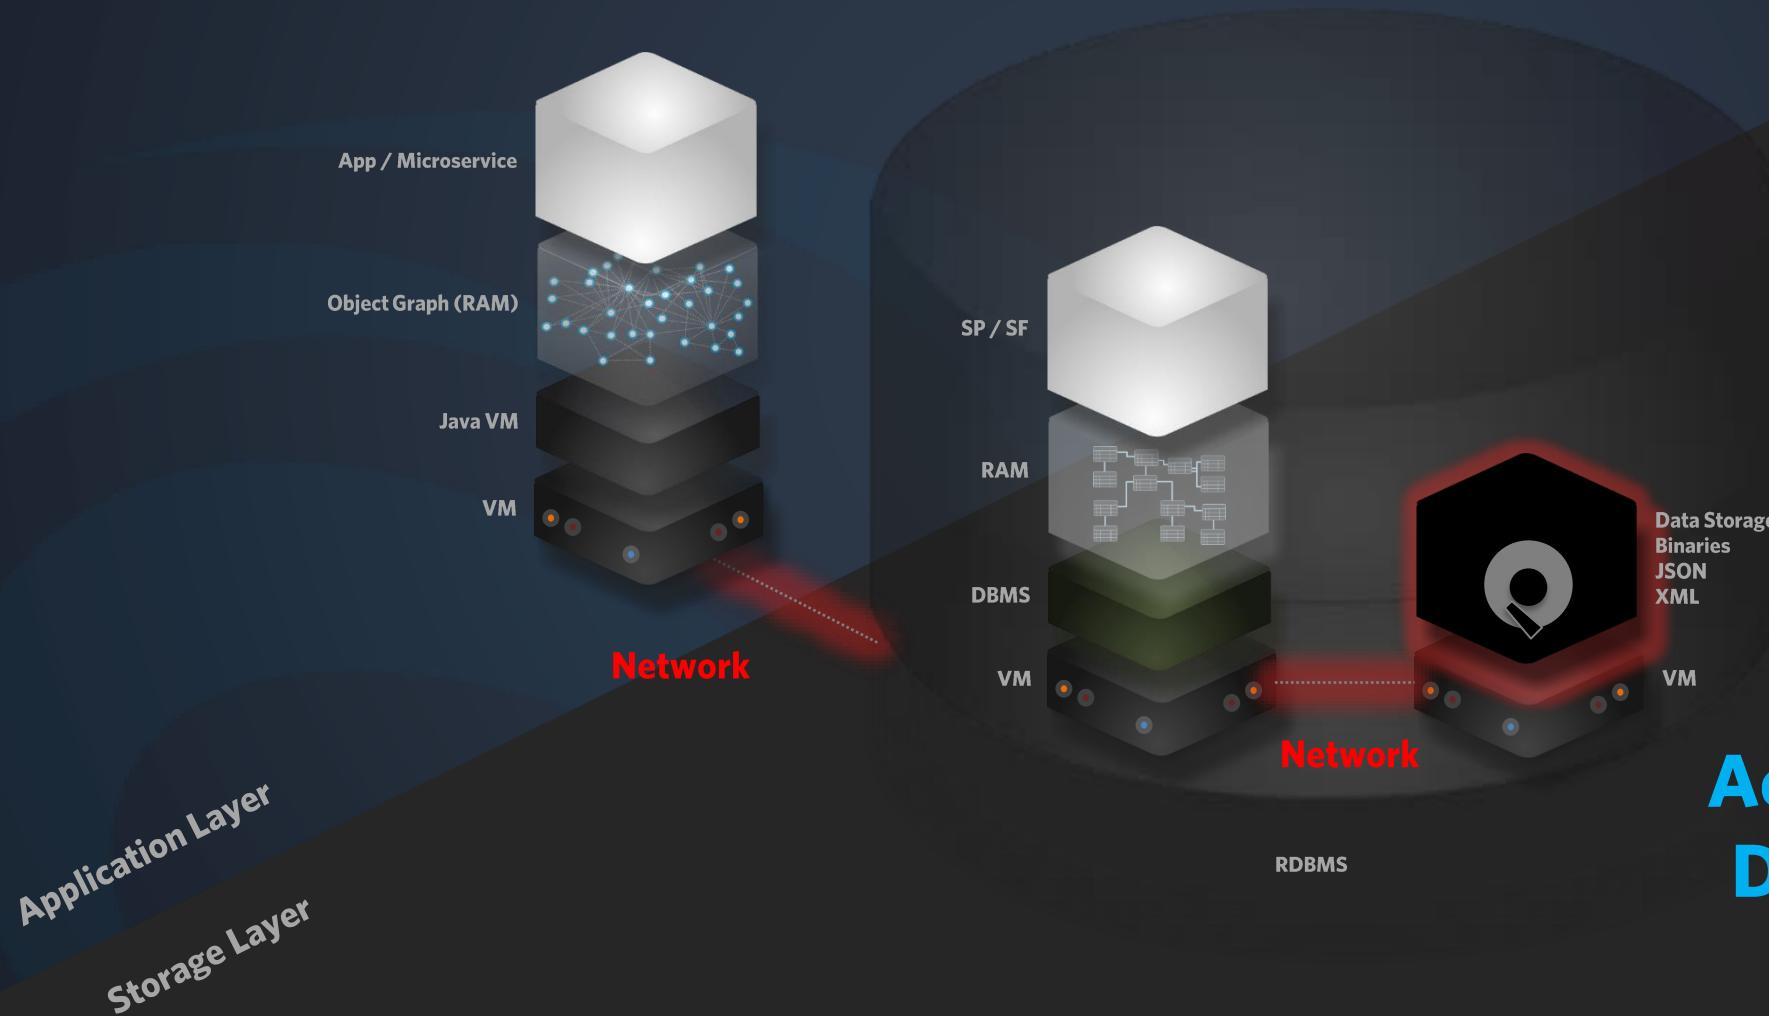
**In-Memory Computing with Java is Extremely Fast:
Microsecond Response-Time!**



Where are the Hidden Latencies?



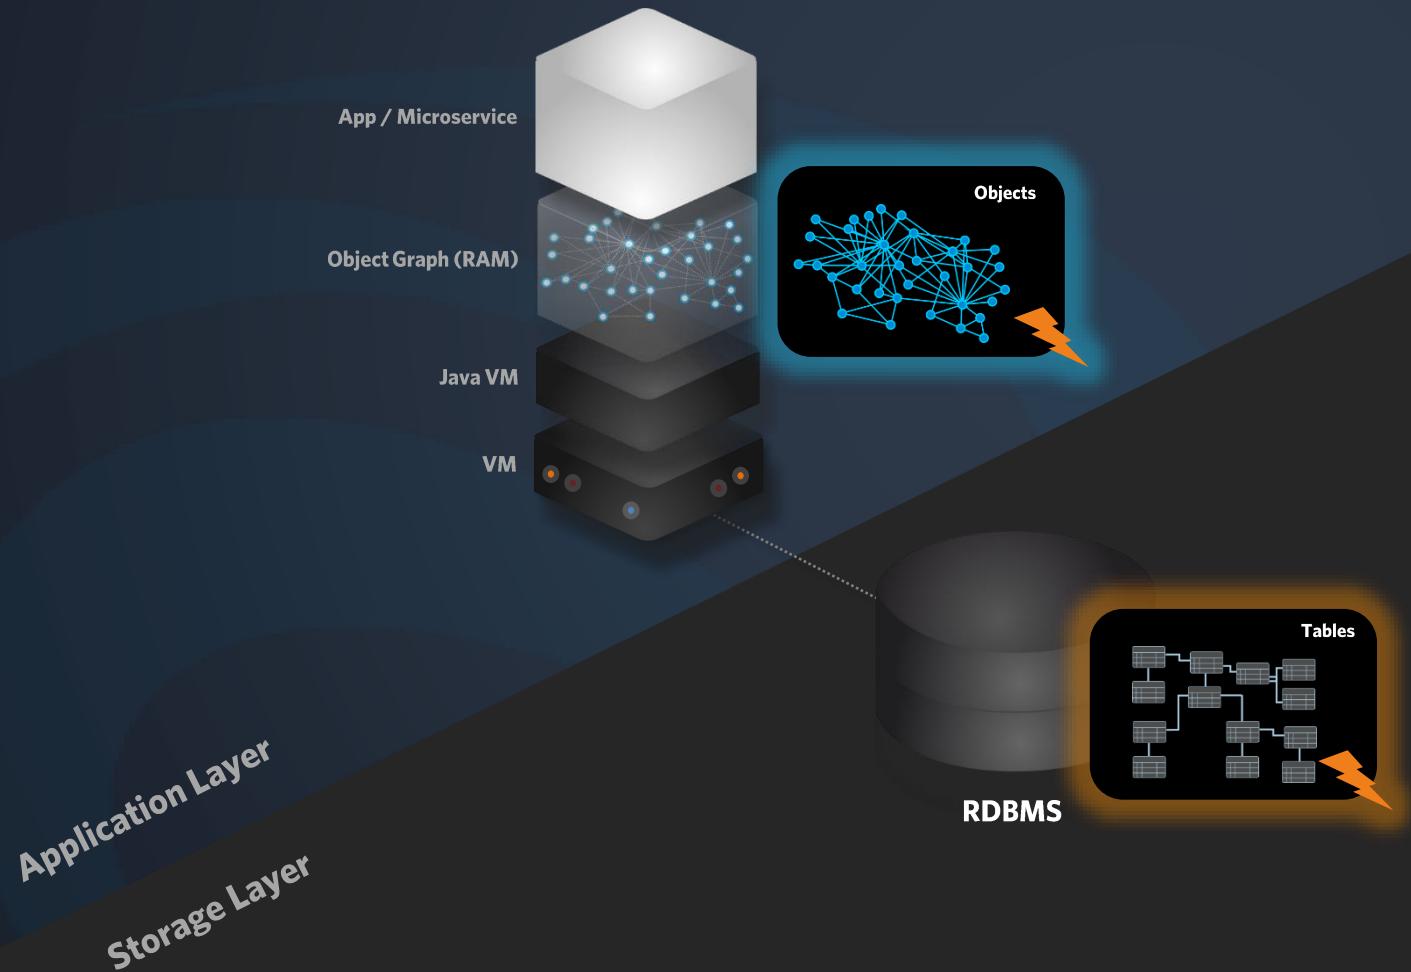
Traditional Database Server in the Cloud



Accessing Hard Disk Drives & Network is Most Expensive.



Java and Relational Databases are Incompatible

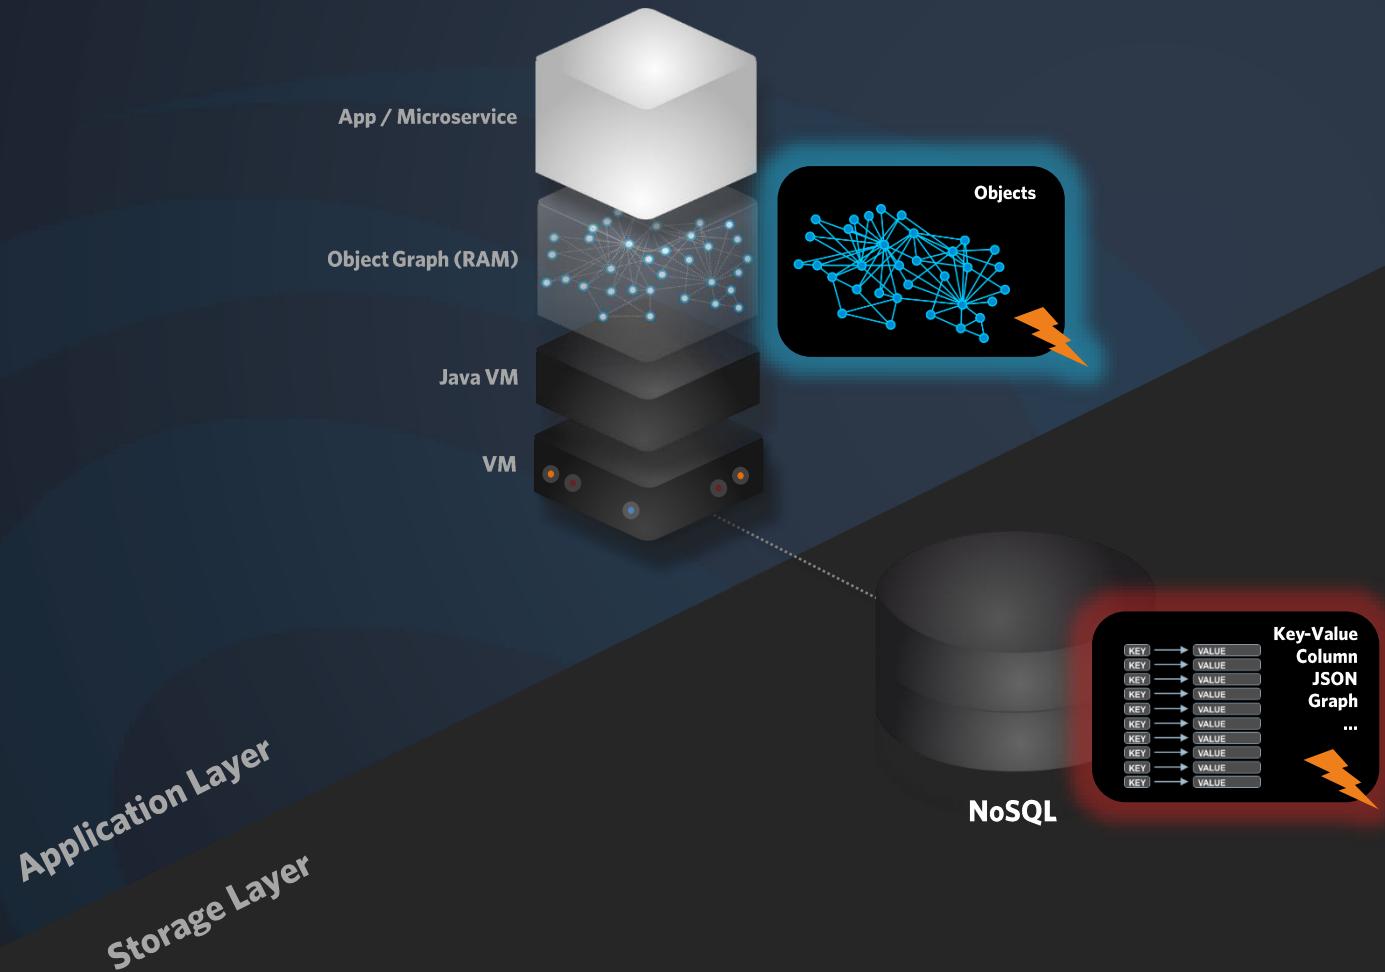


Impedance Mismatches

- **Granularity mismatch**
- **Subtypes mismatch**
- **Identity mismatch**
- **Associations mismatch**
- **Data Navigation mismatch**
- **Data type mismatches**



Java and NoSQL Databases are Also Incompatible



Impedance Mismatches

- **Key-Value**
- **Document** (JSON, XML)
- **Column Store**
- **Graph** (Proprietary)
- **OODB** (Proprietary)
- **Time-Series**

The Problem of Incompatible Data Structures is Well Known as Impedance Mismatch

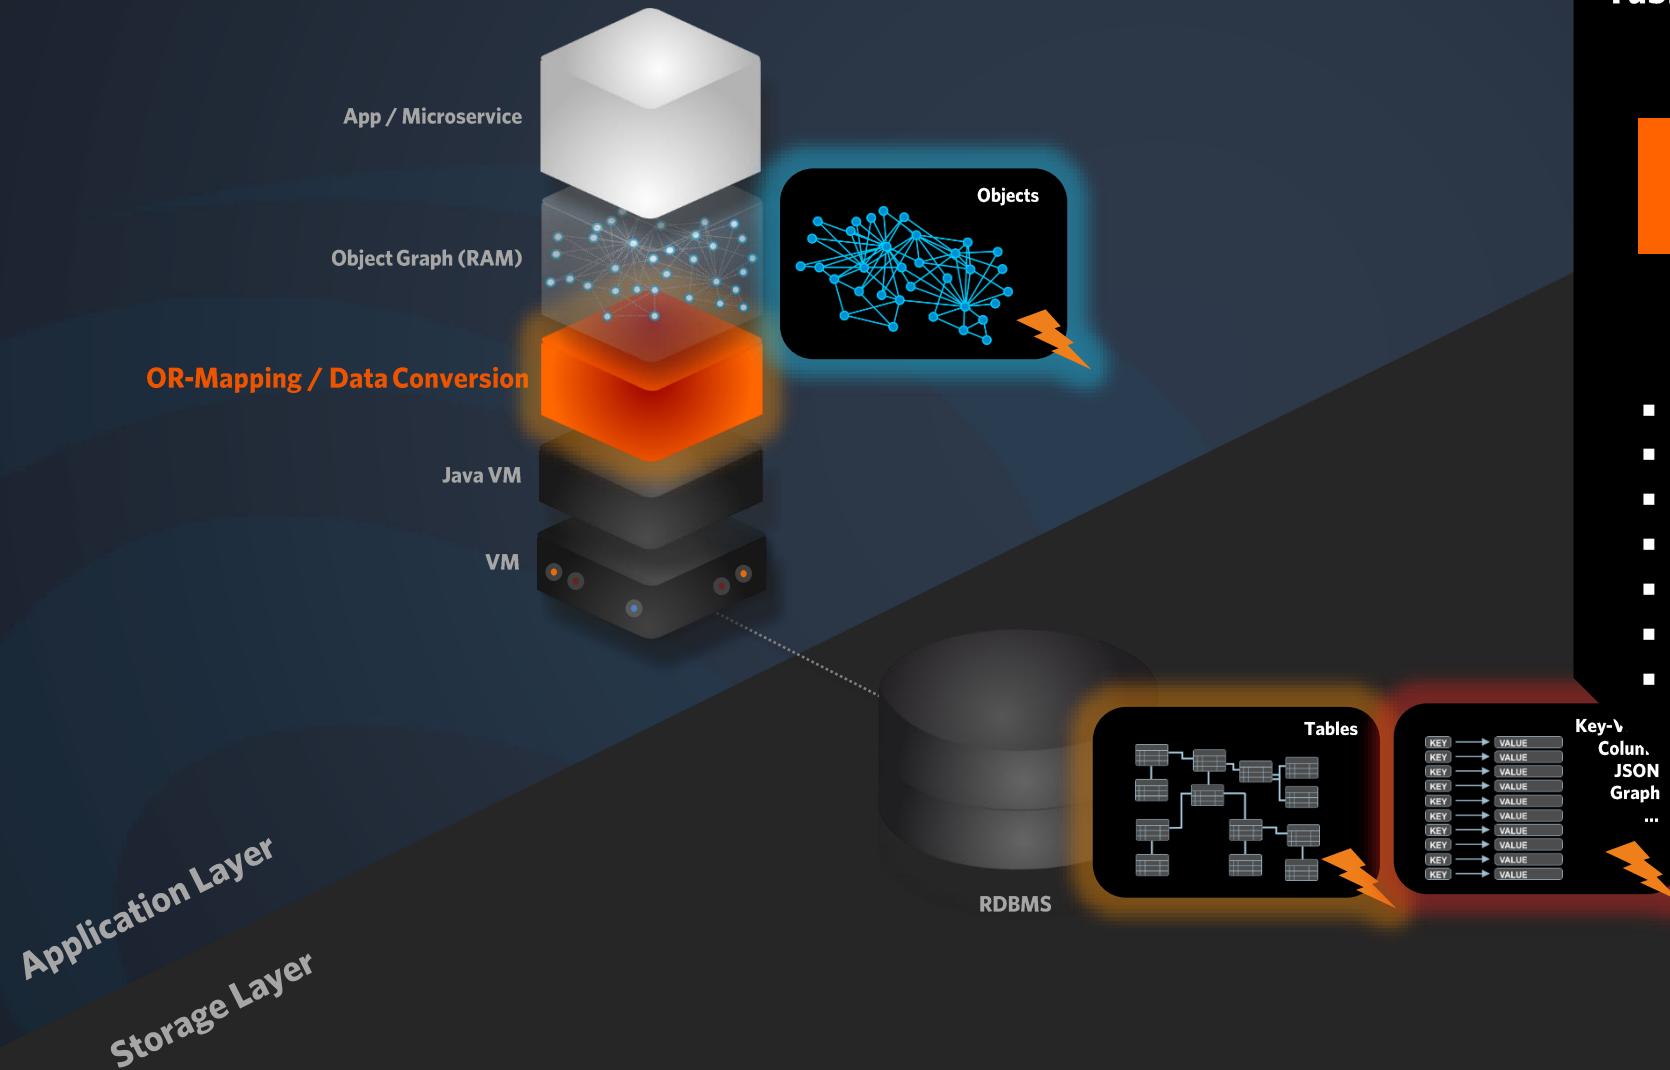


There are various solutions, but they are only a more or less elegant way around the problem. No matter which solution you choose - as long as the systems are different, every developer will sooner or later get to the point where his solution no longer meets one or more of the following points: Maintainability, performance, intelligibility.

”



Object-Relational Mapping / Data Conversion



**Challenge: Storing Objects into
Tables / JSON / Key Value Stores / Graphs**

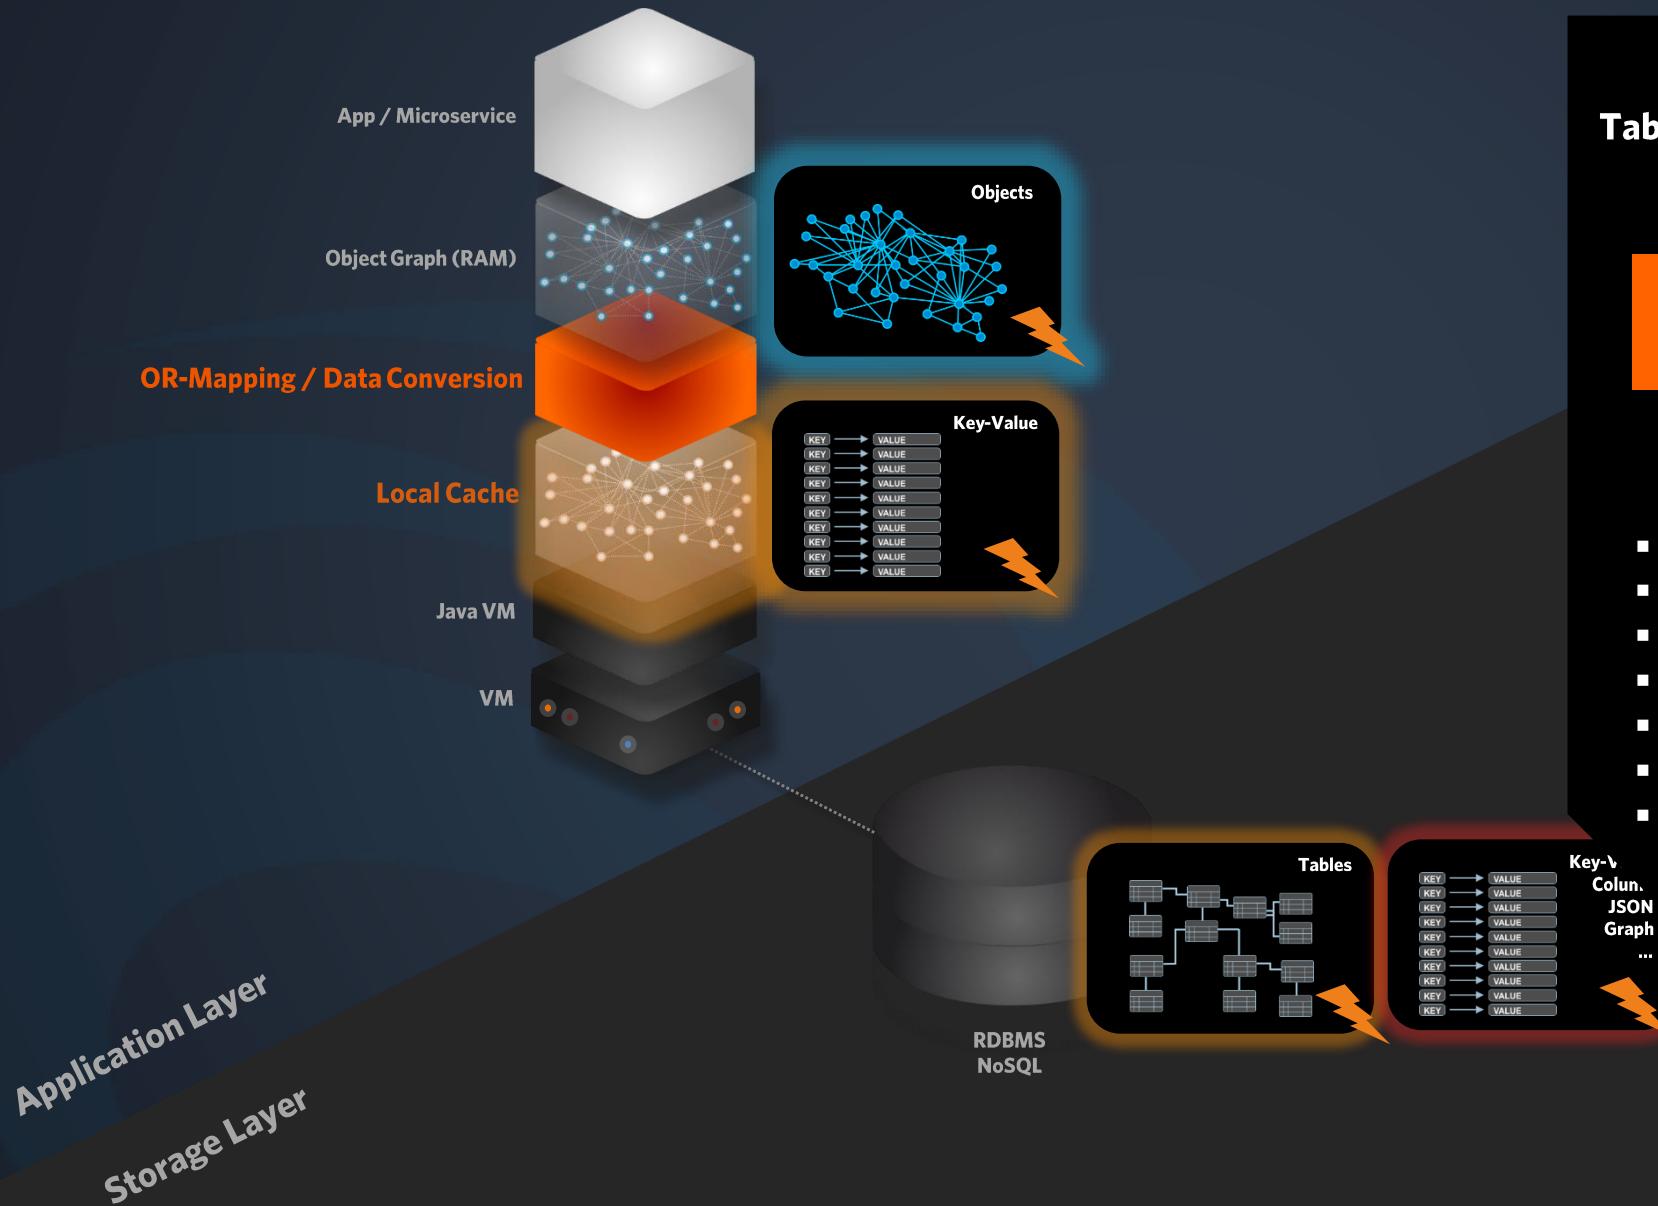
**Data Conversion Through
Every Single Read & Write !**

- Higher CPU consumption
- Higher energy consumption
- Higher CO2 emission
- Higher data center costs
- Expensive latencies
- More complex architecture
- Higher costs of development

**Millisecond
Query Time**



Local Cache: Better Performance, But Additional Complexity



**Challenge: Storing Objects into
Tables / JSON / Key Value Stores / Graphs**

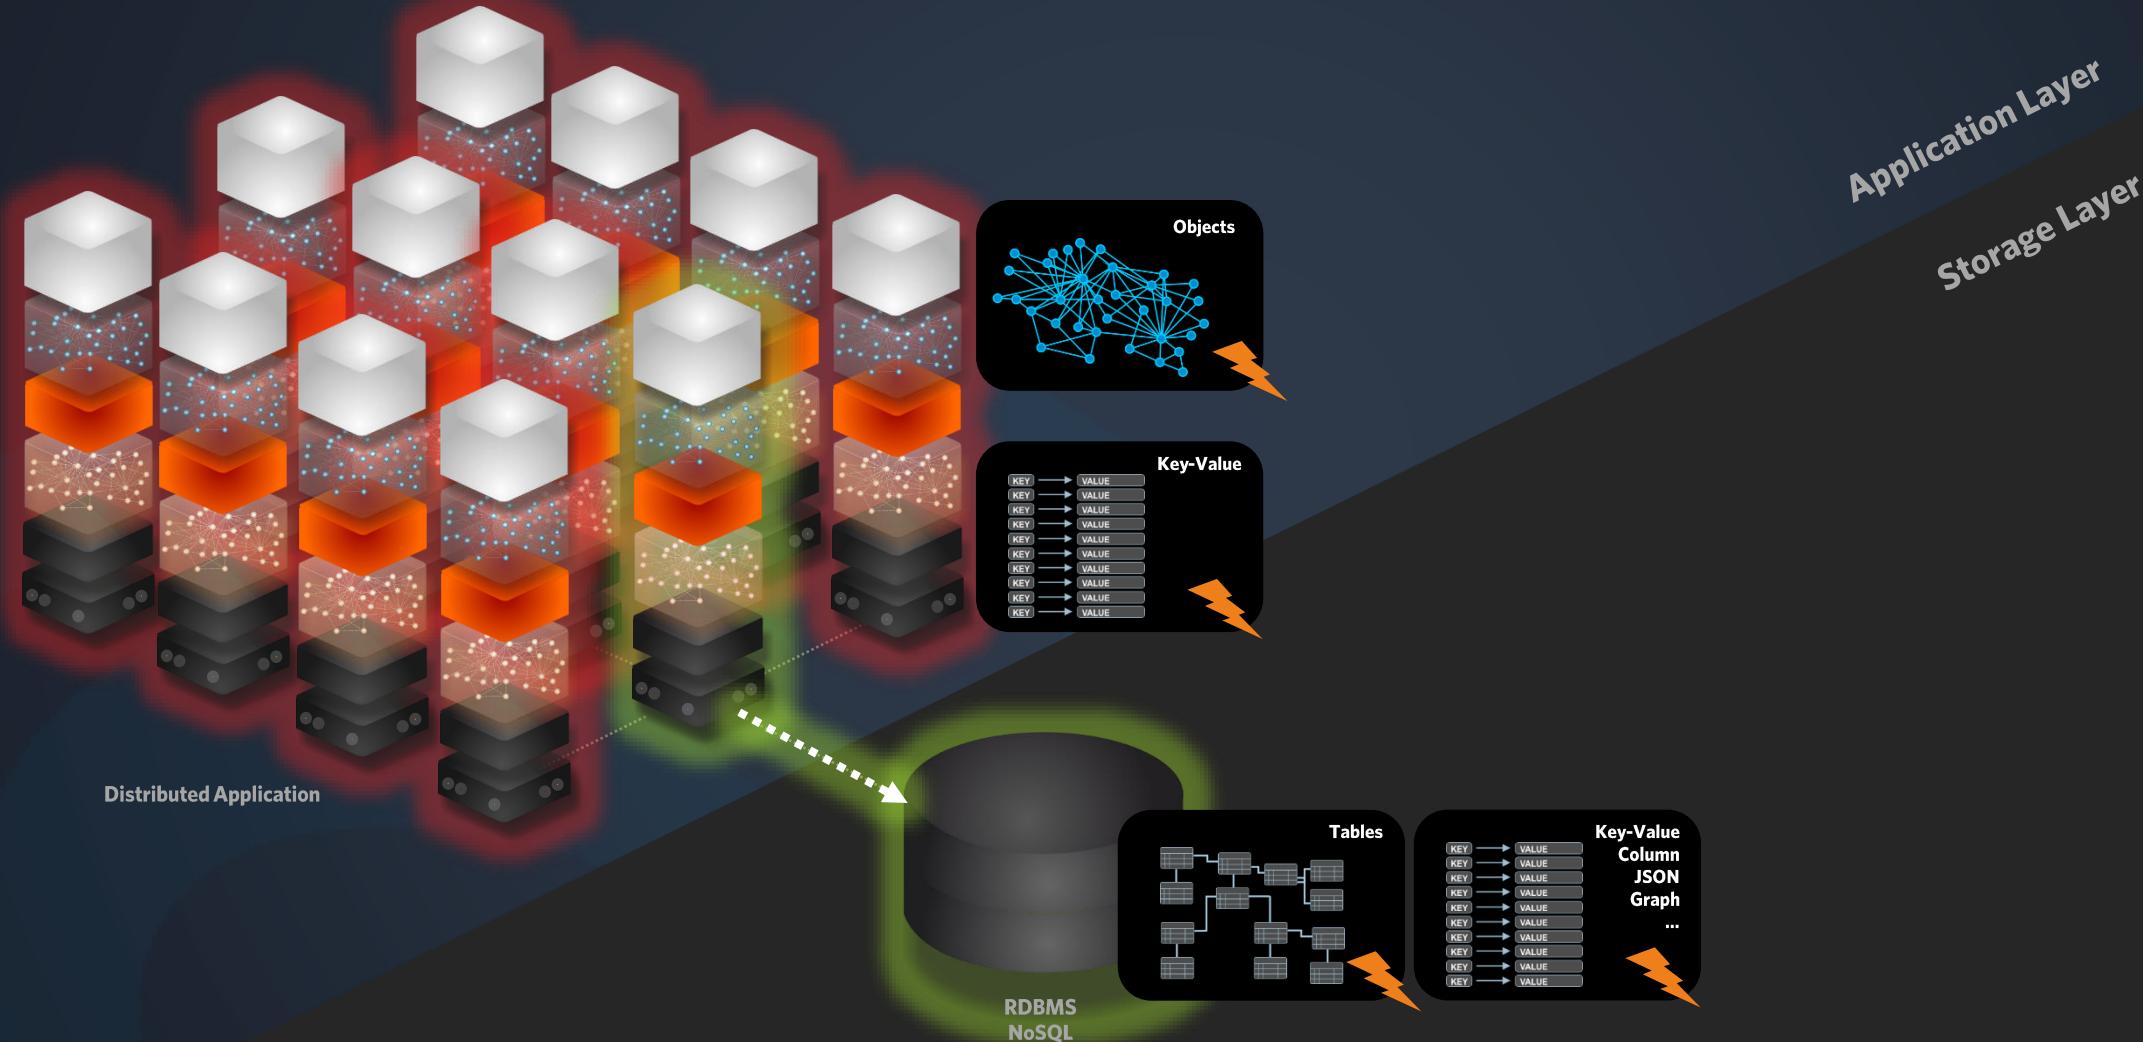
**Data Conversion Through
Every Single Read & Write !**

- Higher CPU consumption
- Higher energy consumption
- Higher CO2 emission
- Higher data center costs
- Expensive latencies
- More complex architecture
- Higher costs of development

**Millisecond
Query Time**

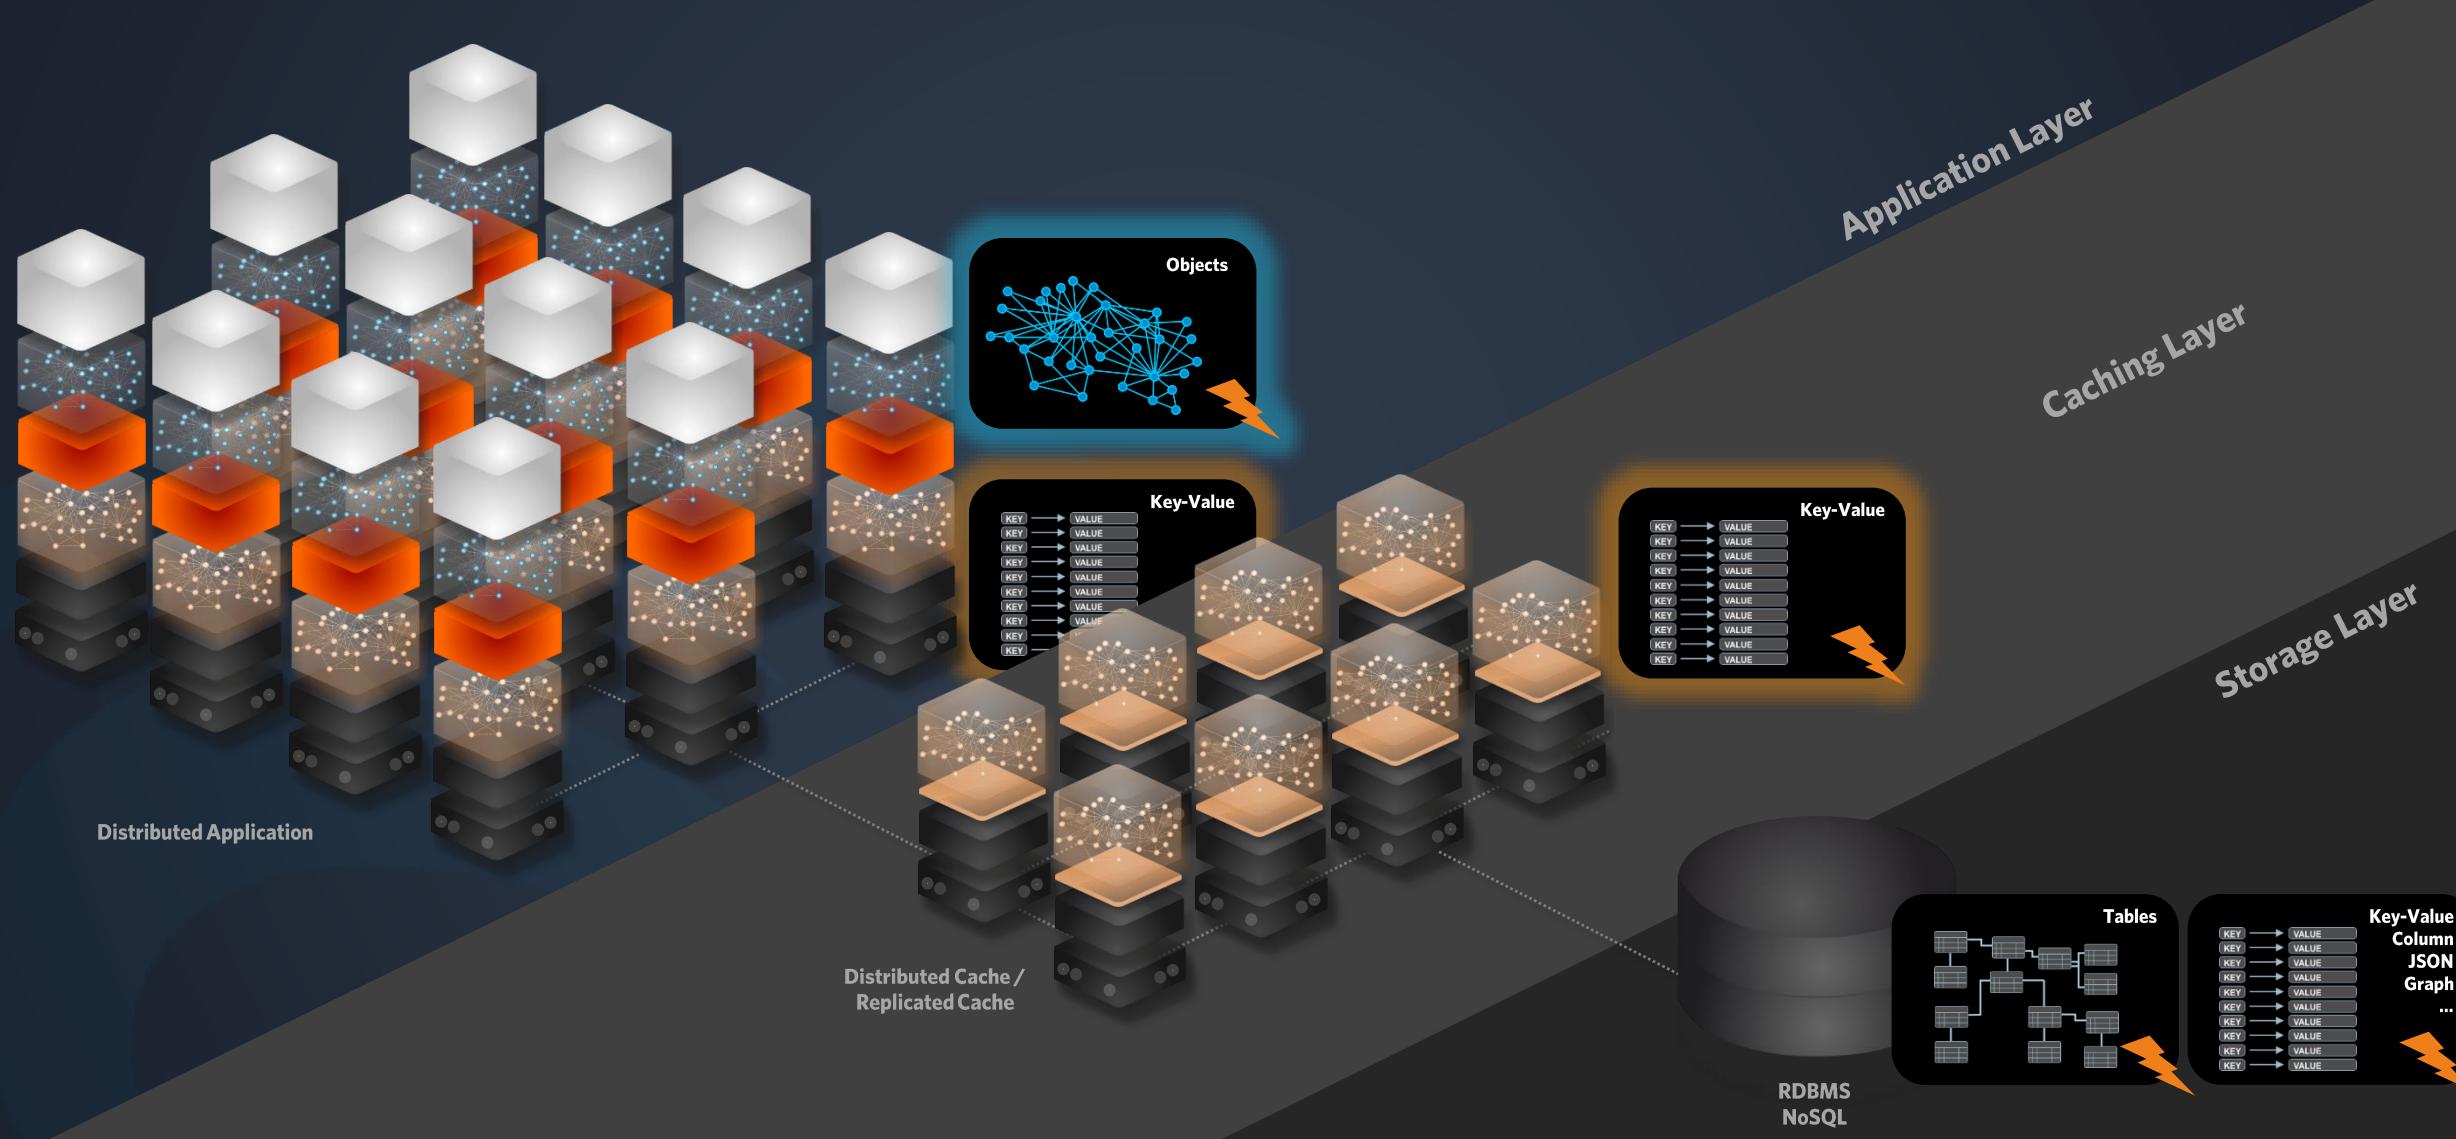


Local Caches in Distributed Apps Can Cause Inconsistent Data



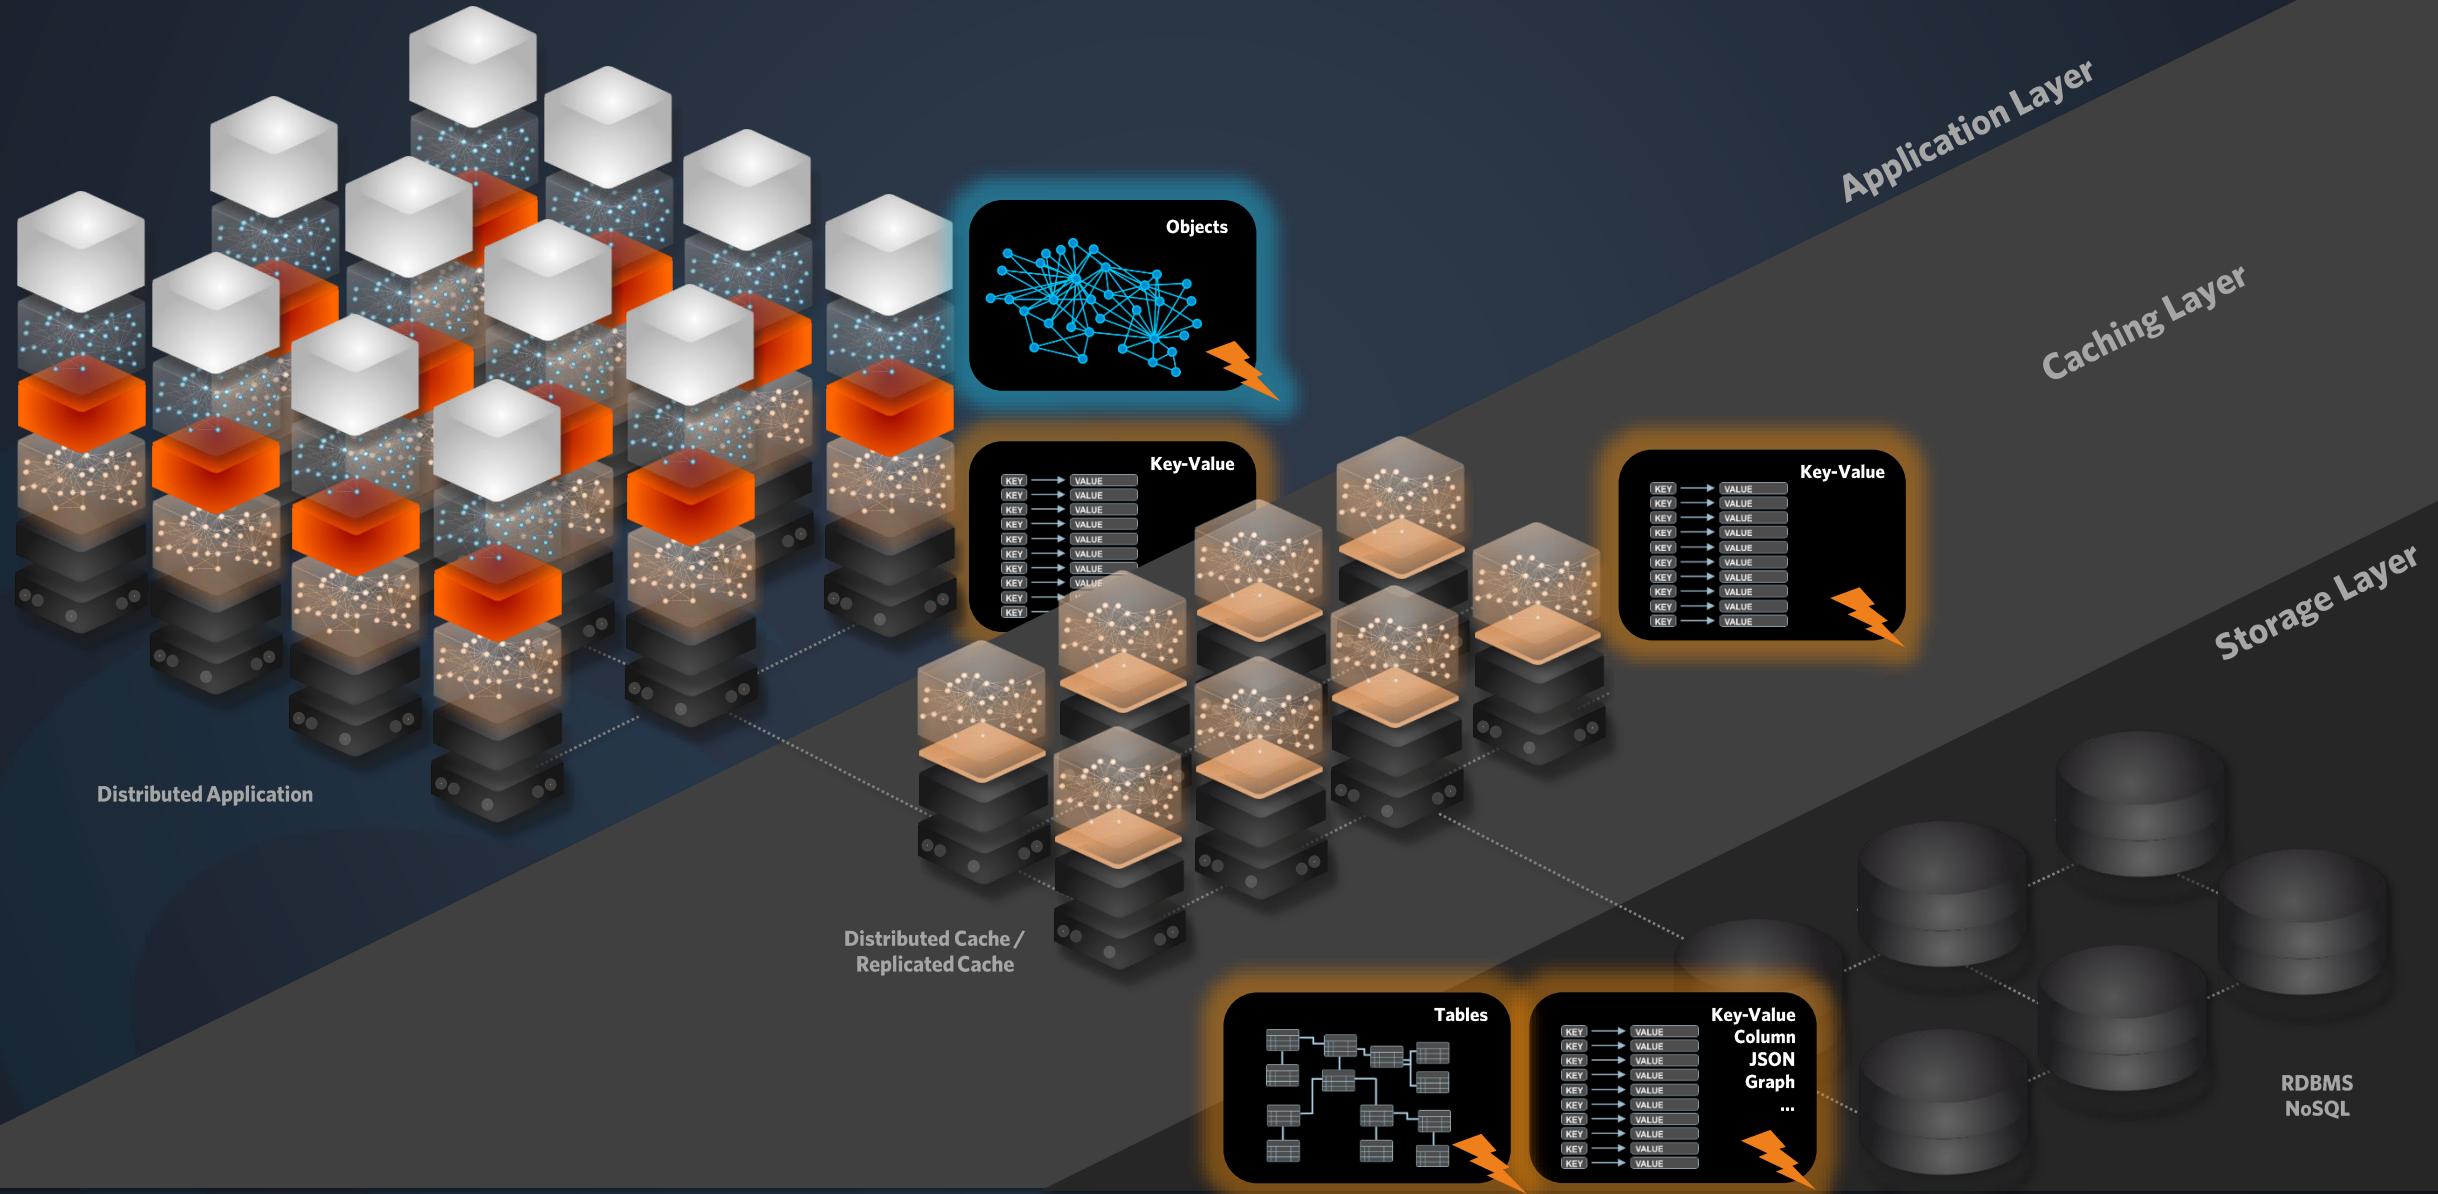


Distributed Cache Architecture



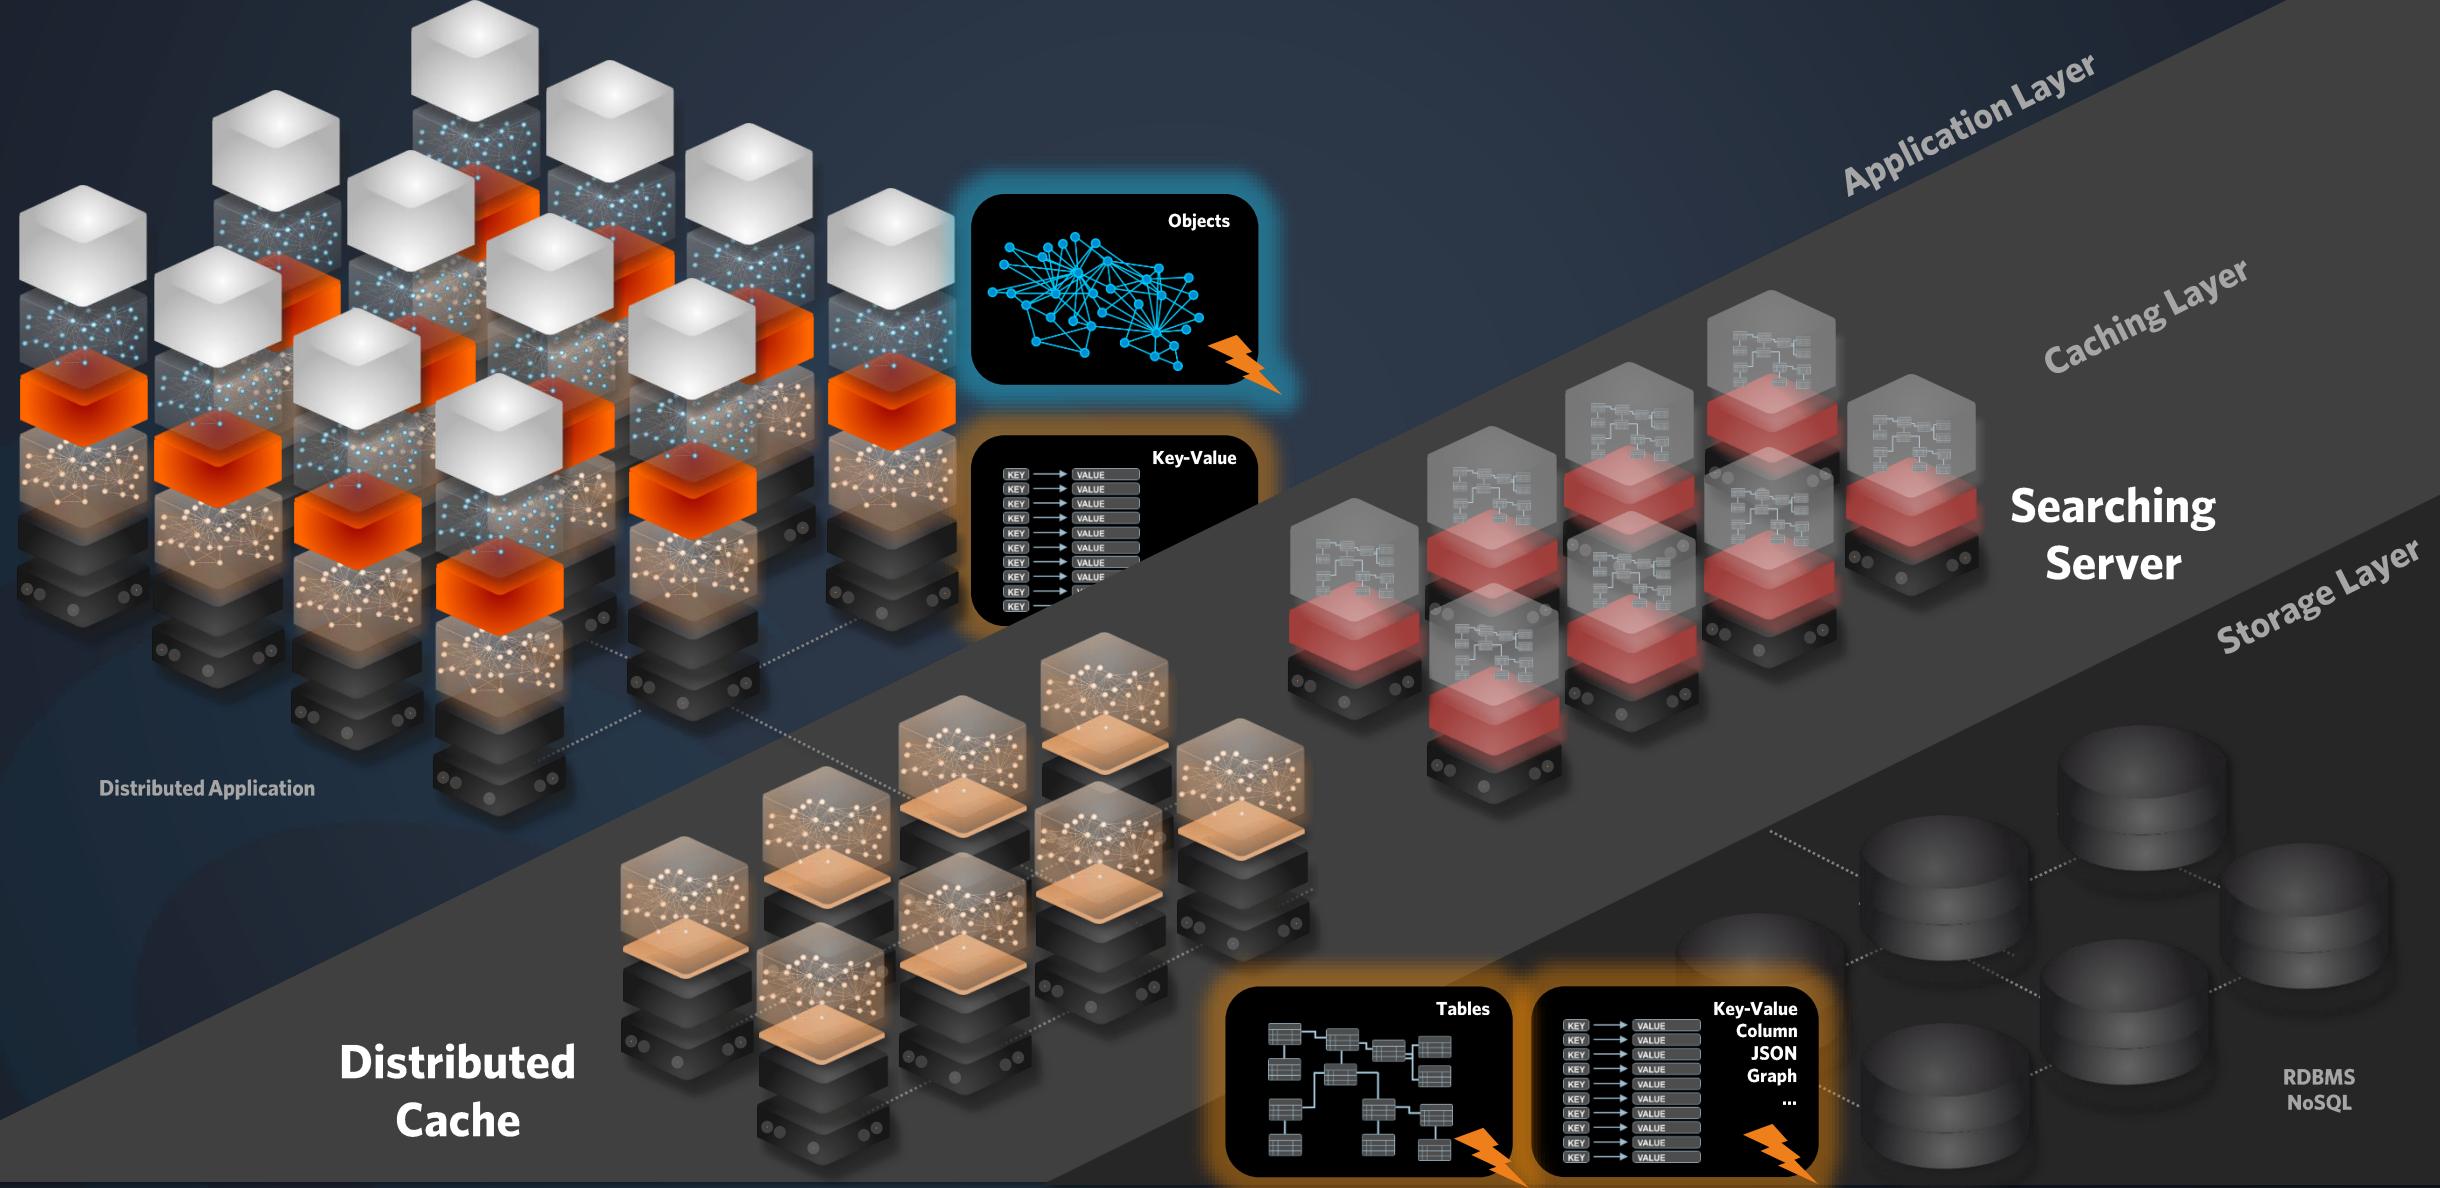


Distributed Systems: 3 Distributed System Layers in Total



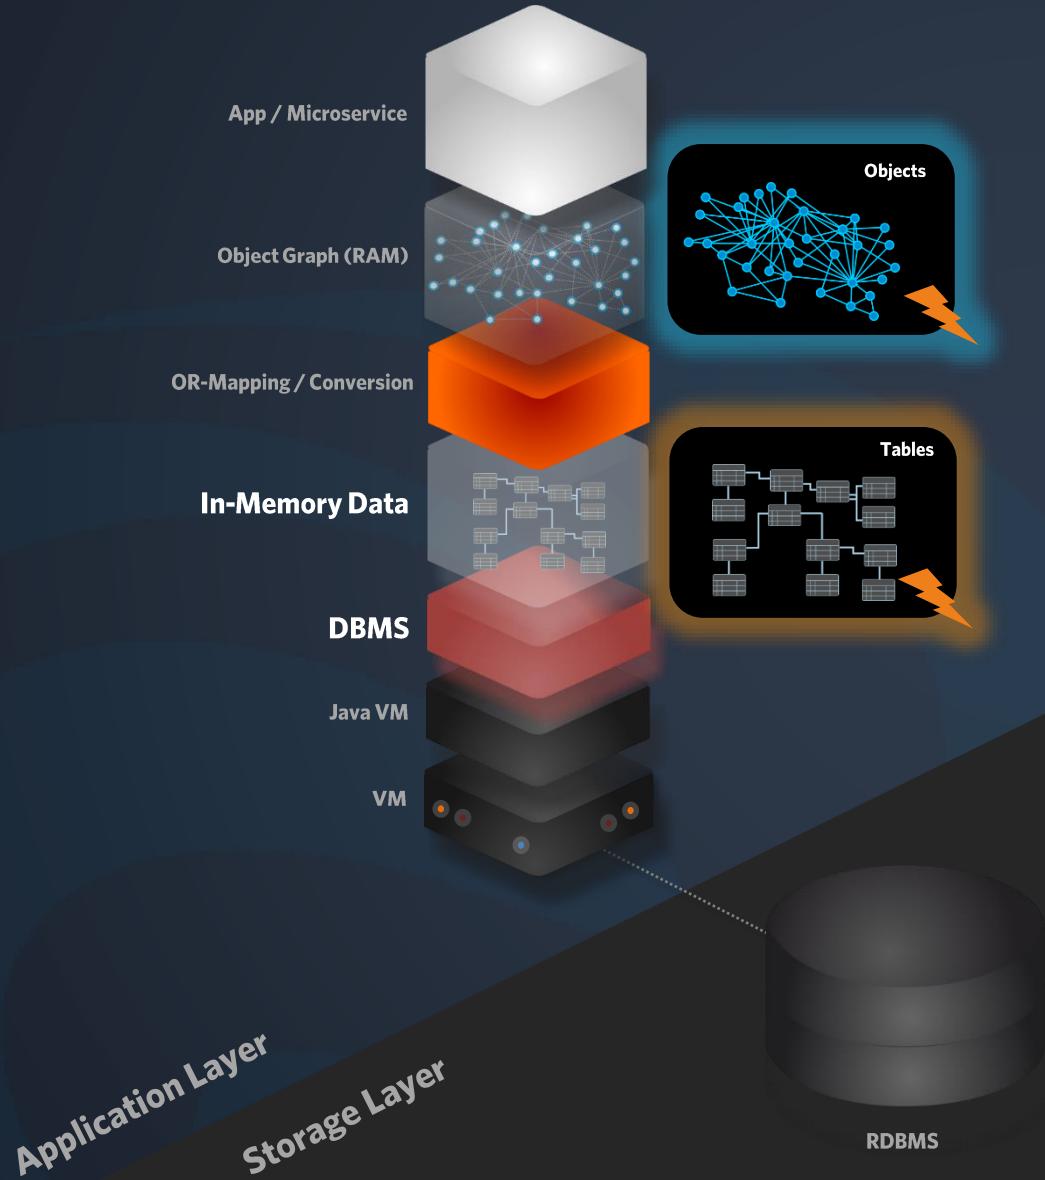


Searching Server



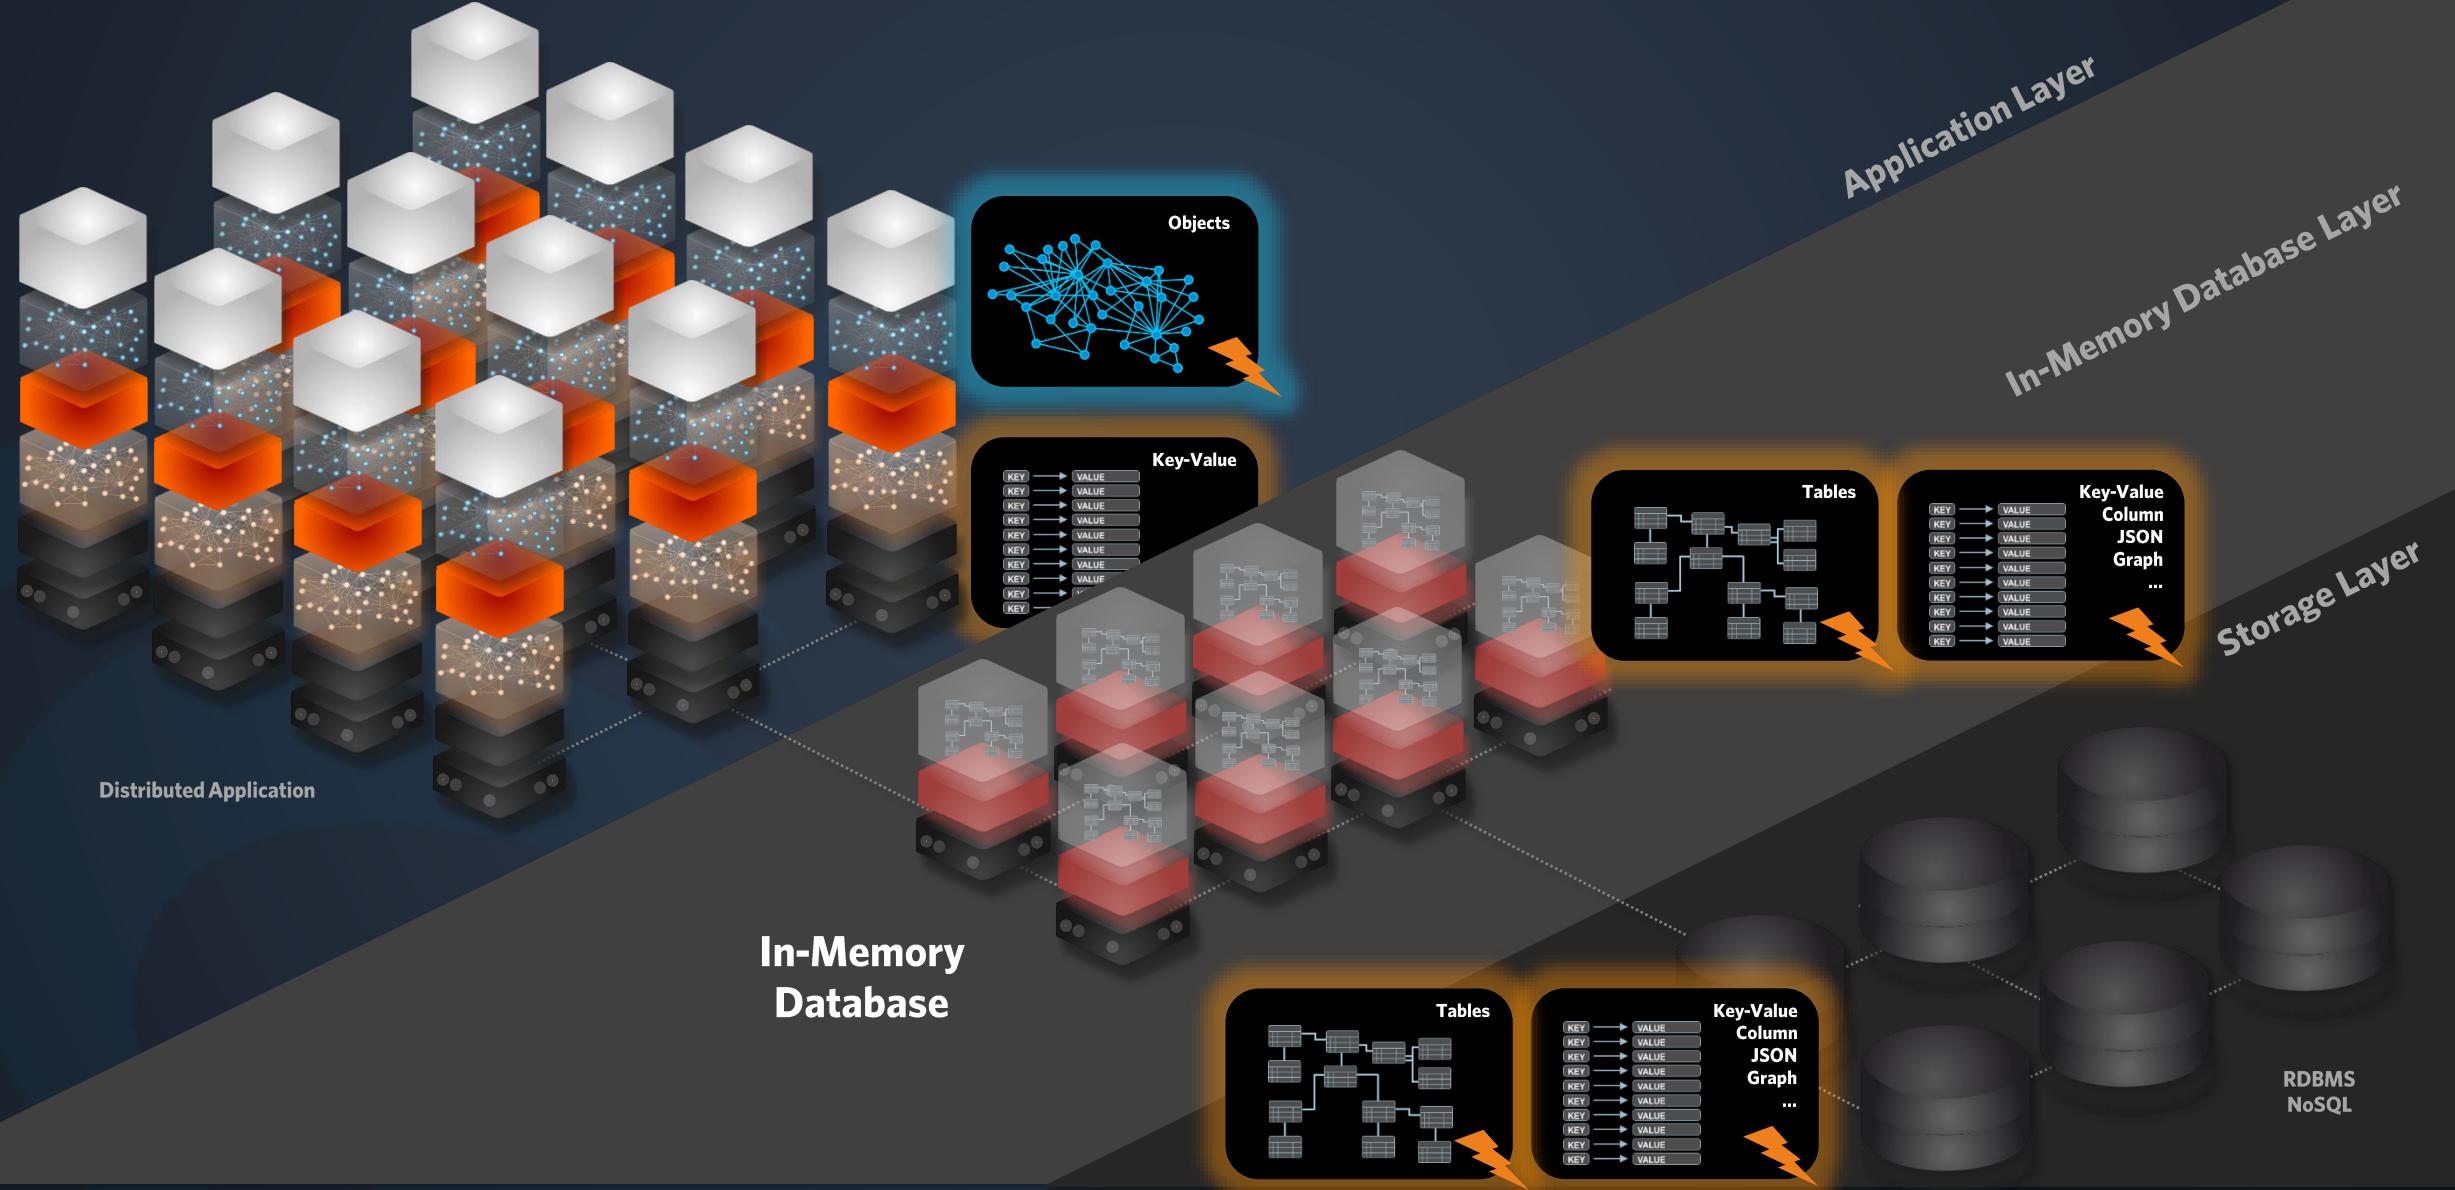


Embedded In-Memory Database Architecture

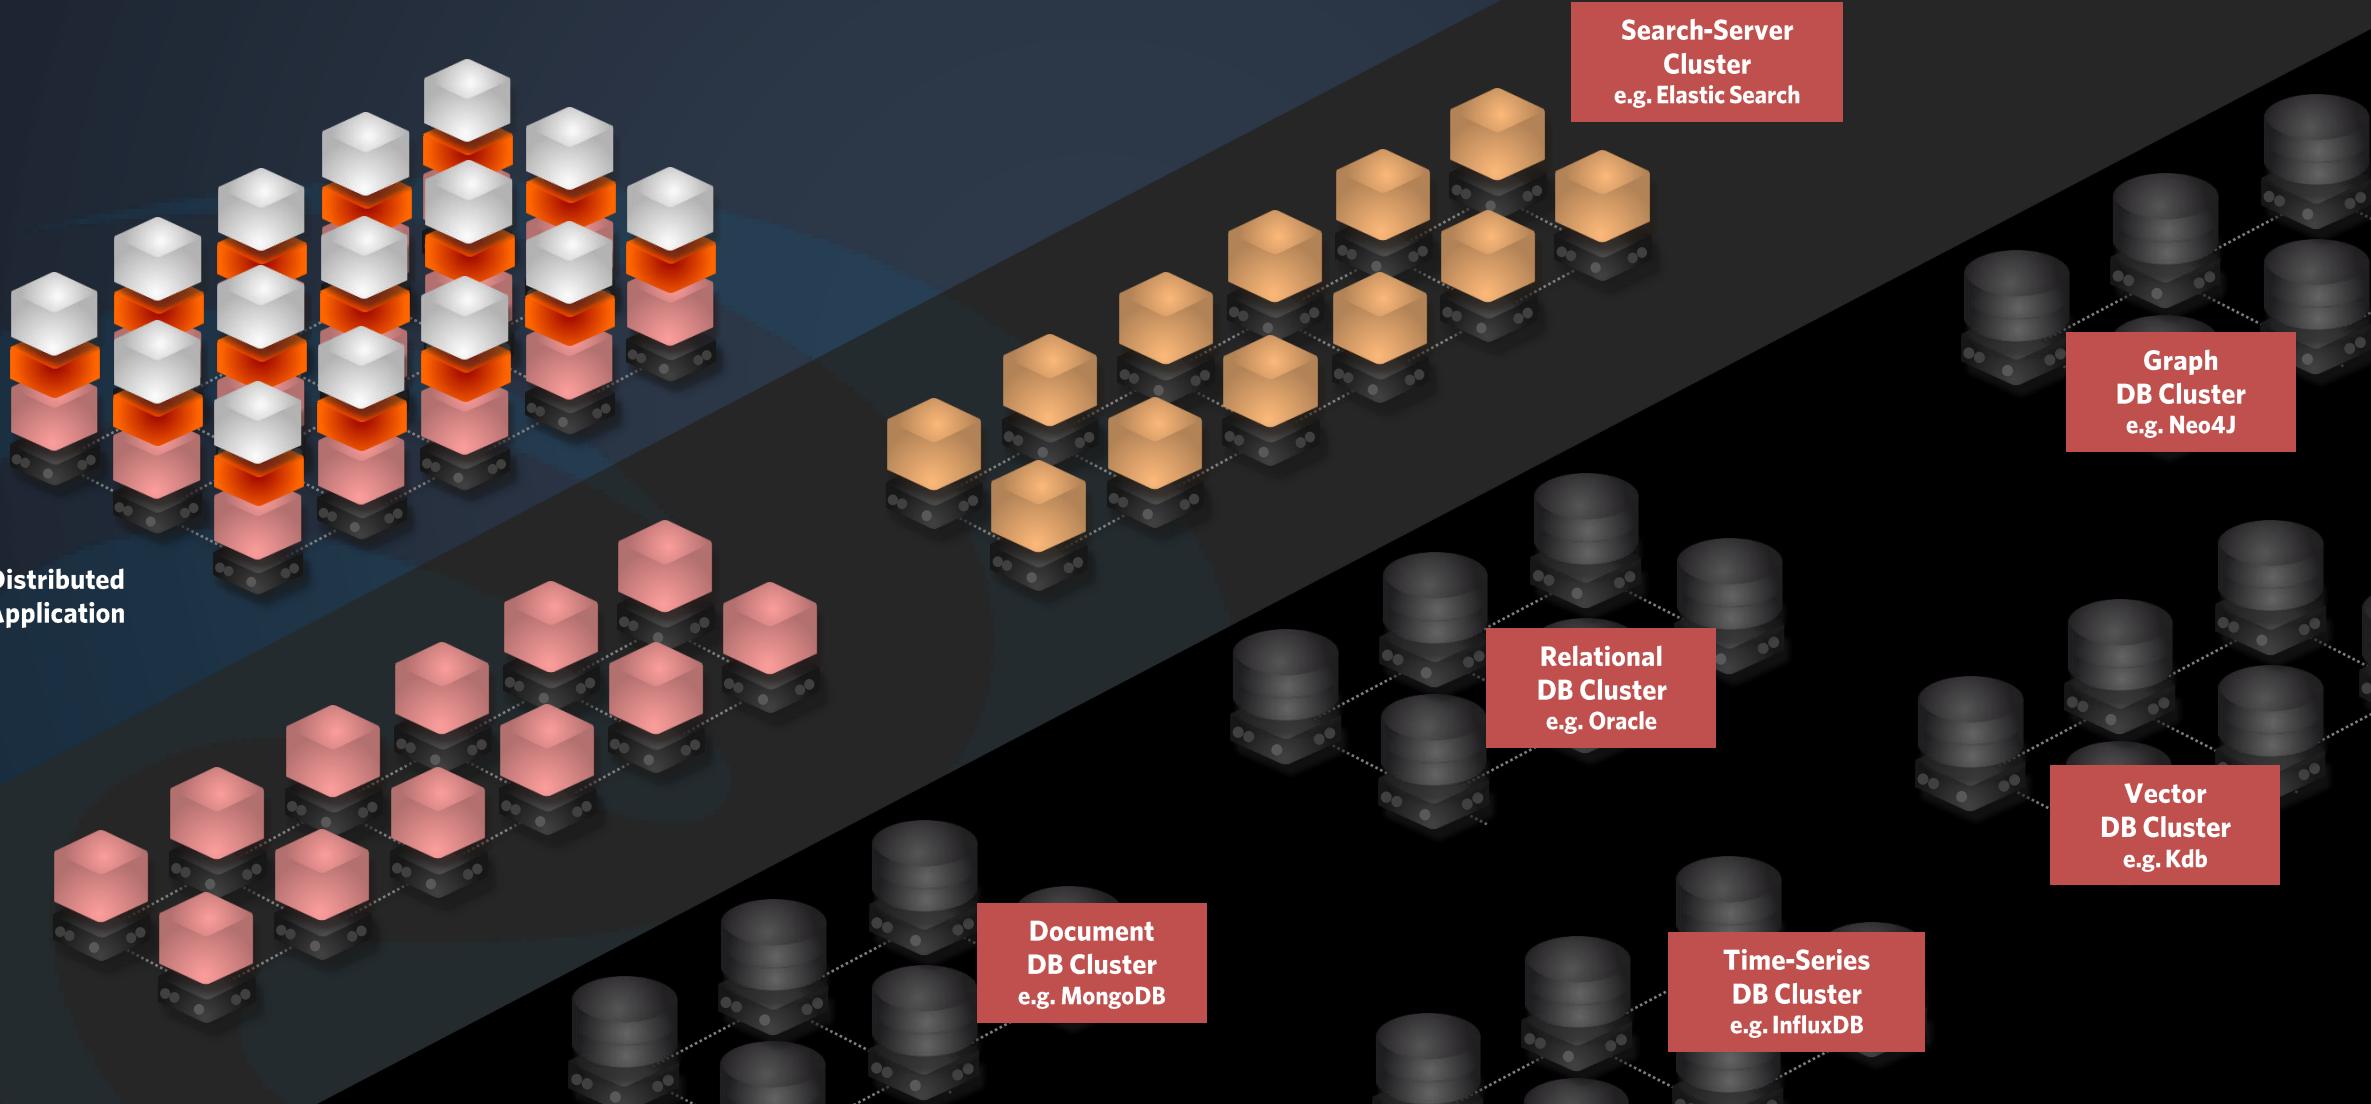




Distributed In-Memory Database Architecture



Today's Enterprise Applications Using Multiple Distributed Solutions. Cloud Costs are Skyrocketing.





Java In-Memory Data Processing



Java is extremely Fast. The JVM's JIT compiler compiles your code into machine code and optimizes it constantly through the runtime. By using newer Java versions, your code will become even better and faster.

Thus, Java can in some cases even outperform C programs which are static. So, everything you build and run with pure Java in the memory will be incredible fast: Data access and searching times in microseconds, realtime responsiveness at pure Java simplicity.



Java



ECLIPSE™ FOUNDATION OpenJ9

Java's Object Graph Model is a Multi-Model Data Structure

Any Data Model

- Objects
- Collections
- Graphs
- Documents (JSON)
- Vectors



```
public static void booksByAuthor()
{
    final Map<Author, List<Book>> booksByAuthor =
        ReadMeCorp.data().books().stream()
            .collect(groupingBy(book -> book.author()));

    booksByAuthor.entrySet().forEach(e -> {
        System.out.println(e.getKey().name());
        e.getValue().forEach(book -> {
            System.out.print('\t');
            System.out.println(book.title());
        });
    });
}
```



Searching & Filtering with Java Streams API

1000x

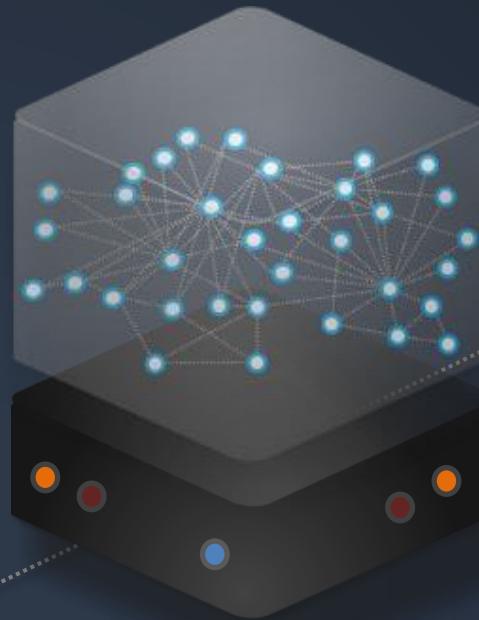
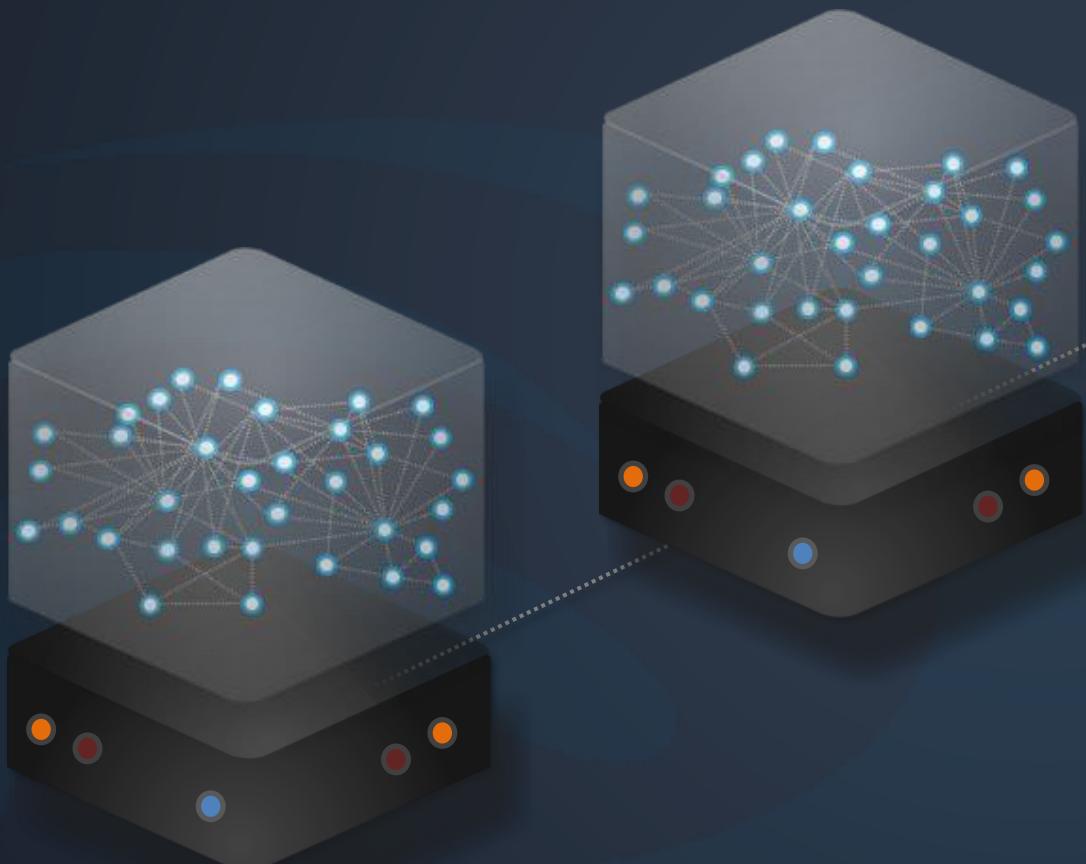
Databases: **Millisecond Query Time**
Java Streams API: **Microsecond Query Time**





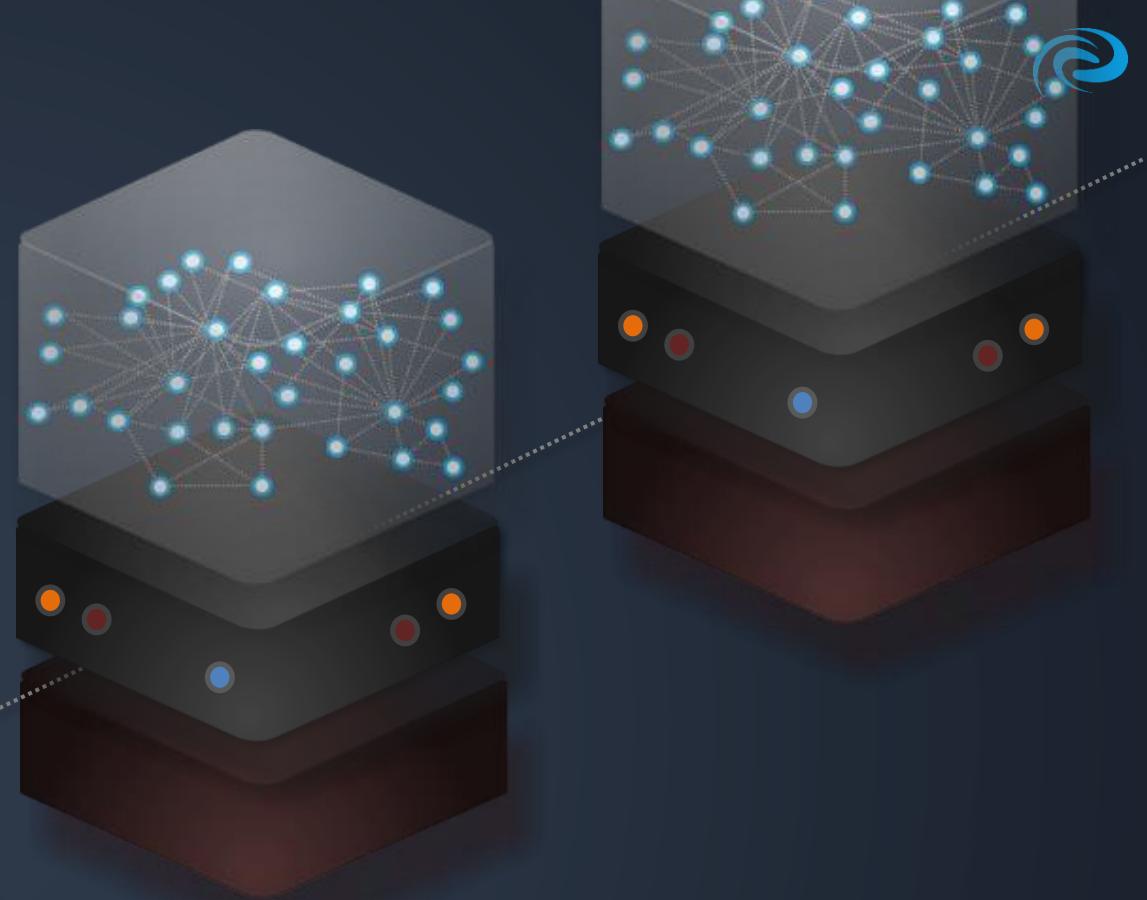
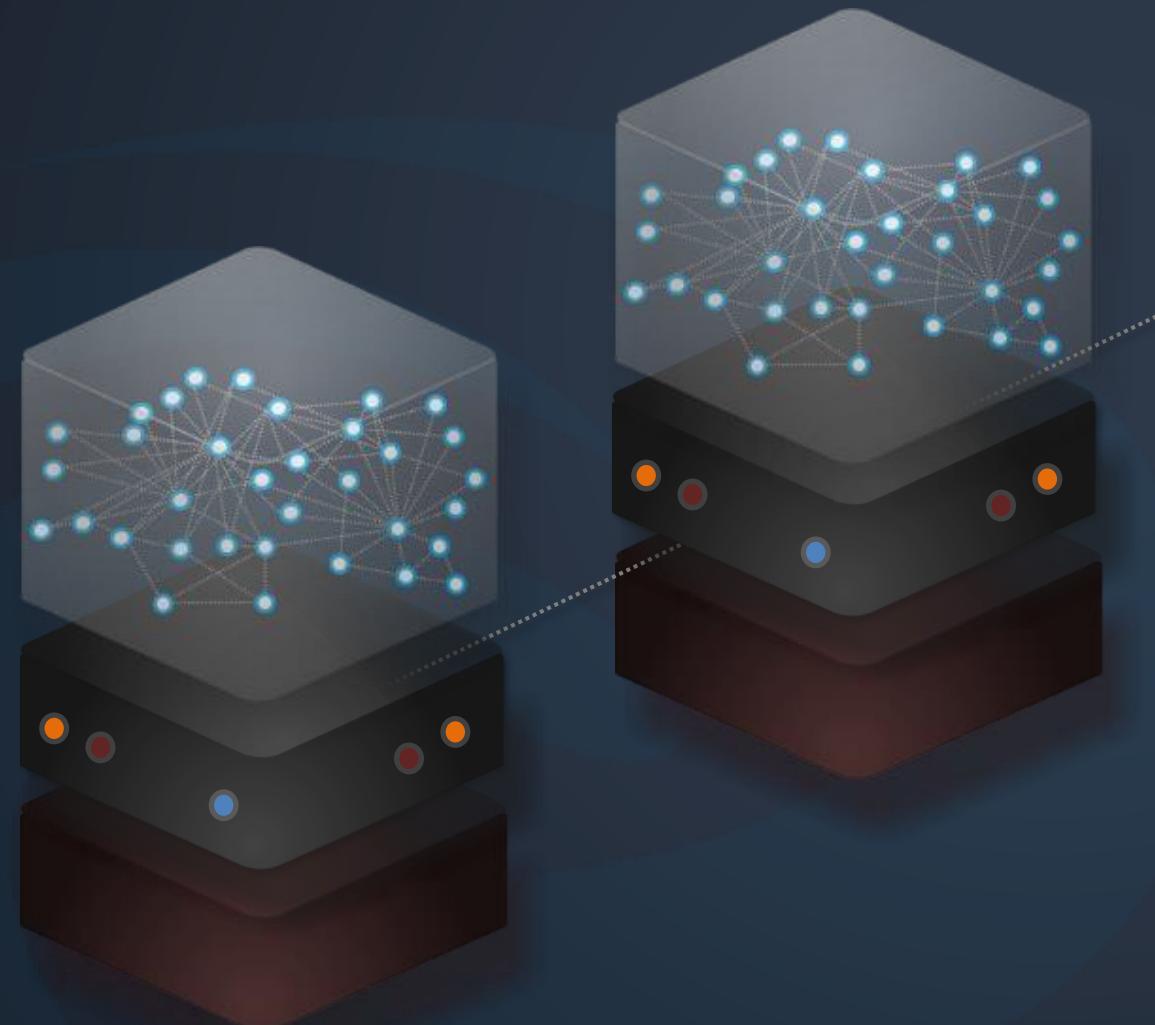
Missing Features in Java For In-Memory Data Processing

Missing Features in Java Enabling In-Memory Data Processing:



**Java Object Graphs
Cannot Be Distributed
and Replicated.**

Missing Features in Java Enabling In-Memory Data Processing:



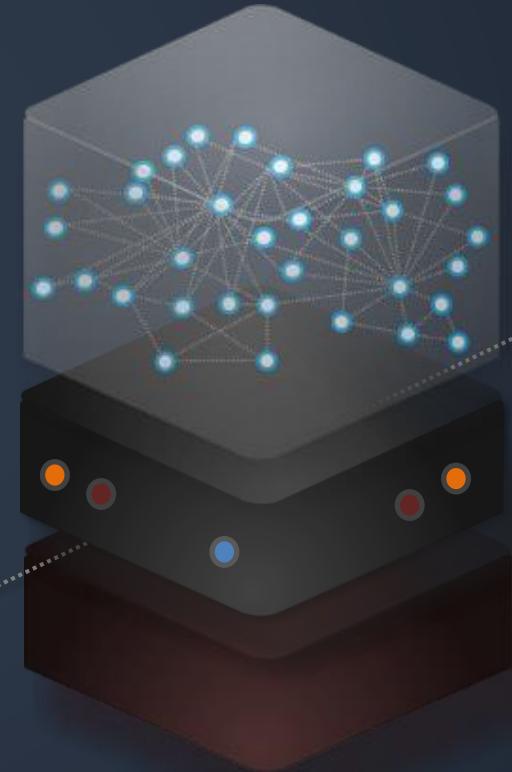
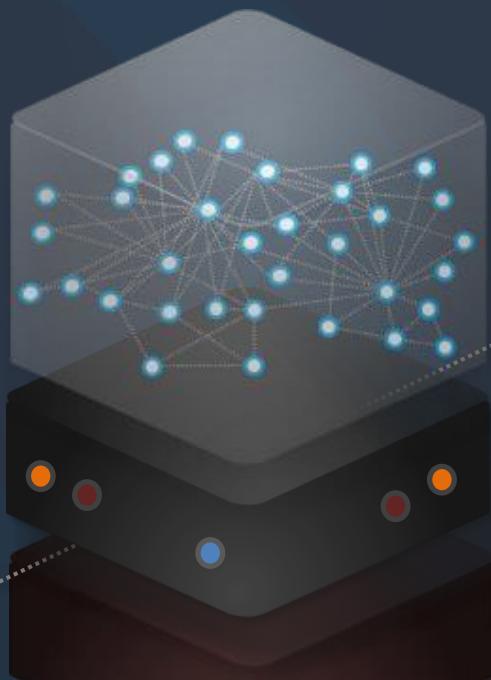
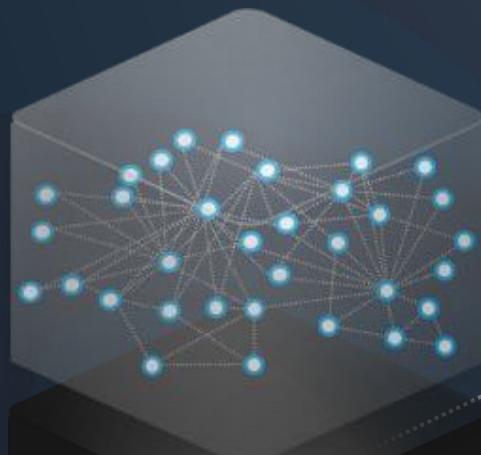
**State Cannot Be Persisted
Highly Efficient. Just
Snapshots are Not Efficient.**



Missing Features in Java Enabling In-Memory Data Processing:

- **Distributed Java Object Graphs**
- **State Persistence**
- **ACID Transactions**
- **Indexing**
- **Schema Migration**

and more





Imagine You Get All of These Features ...



Eclipse Data Grid

Ultra-Fast In-Memory Data Processing with Pure Java.

Caching. Searching. Indexing. Java Business Logic. Persistence. Distributed. Open-Source.

Java In-Memory Data Processing Foundation.



Eclipse Data Grid Features

Indexing
Extreme Fast Off-Heap Bitmap Index to Minimize IO Operations.

GigaMap
Powerful Map Enabling Fully Automated Lazy-Loading.

Locking API
Enabling Locking and Simplifying Concurrency Handling on the Java Object Graph Level.

Lucene Integration
Extends Eclipse Data Grid With Full-Text Search Capabilities.

ACID Persistence
Persisting Any Java Object Graph Transaction-Safe In Various Storage Targets.

Highly-Secure Serialization
Persisting Any Java Object Graph Transaction-Safe In Various Storage Targets.

Object Graph Replication
Distributing Any Java Object Graph Through Multiple JVMs in a Cluster.

Cluster
Configuration Files For Creating an Eclipse Data Grid Cluster Environment Based on Kubernetes.

Native Java Features

Native Java Object Model
Eclipse Data Grid Supports the Native Java Object Model. No Mappings Or Conversions. Just POJOs.

Java Streams API
Searching Java Object Graphs with Java Streams API.

Any Logic in Java
Implement Any Business Logic in Java.

OpenJ9 JVM
OpenJ9 Is Highly Performant and Memory Efficient, and Provides Fast Start-Up Time.





Step 1: Java In-Memory Computing: Pure Java Concepts Only.

**Forget database concepts,
focus on Java concepts
only! Use Core Java
concepts and any Java
library and framework.**

- Design any complex Java object model
- Use any Java types
- Use just POJOs
- Search with Java Streams

...



Business Logic

Any Java Object Graph

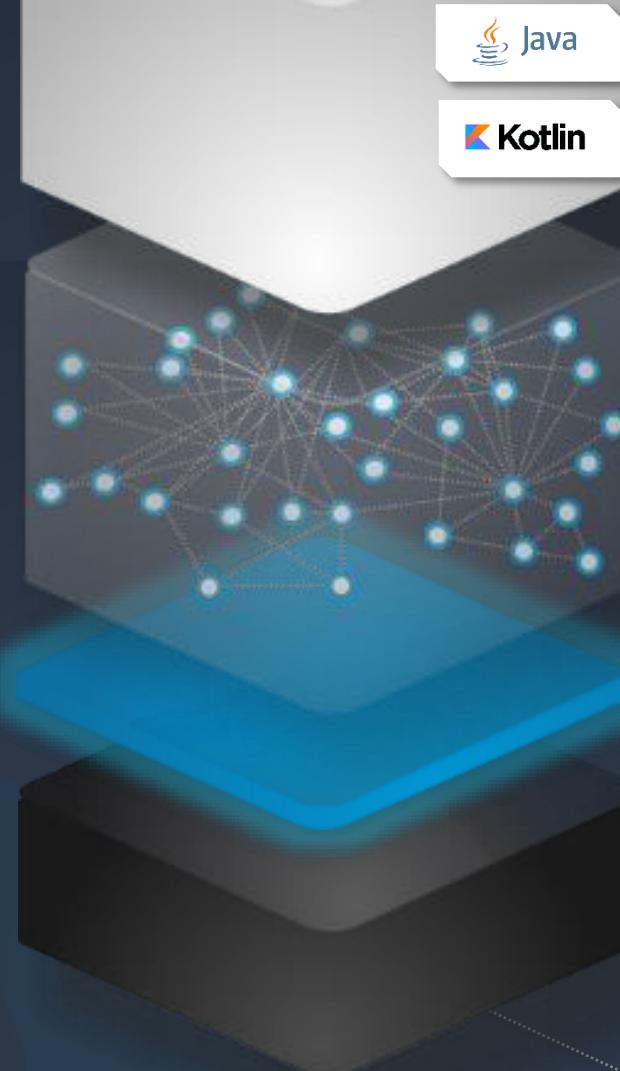
JVM



Step 2: Built-In Persistence and Replications with EclipseStore

Eclipse**Store** 

EclipseStore is a persistence engine to store any Java object graph of any size and complexity into any binary data storage, and to restore it completely or partially in RAM on demand.



Any Java Object Graph

www.eclipsestore.io

Any Data Storage
(Binaries)

0 1 0 0 0 1 0
0 0 0 0 0 0 0
1 0 1 1 1 0 1
0 1 0 0 0 1 0
0 1 0 0 0 1 0

Step 3: Deploy & Run Your App with Eclipse Data Grid as a Cluster

**Queries in Java
Any Business Logic in Java**

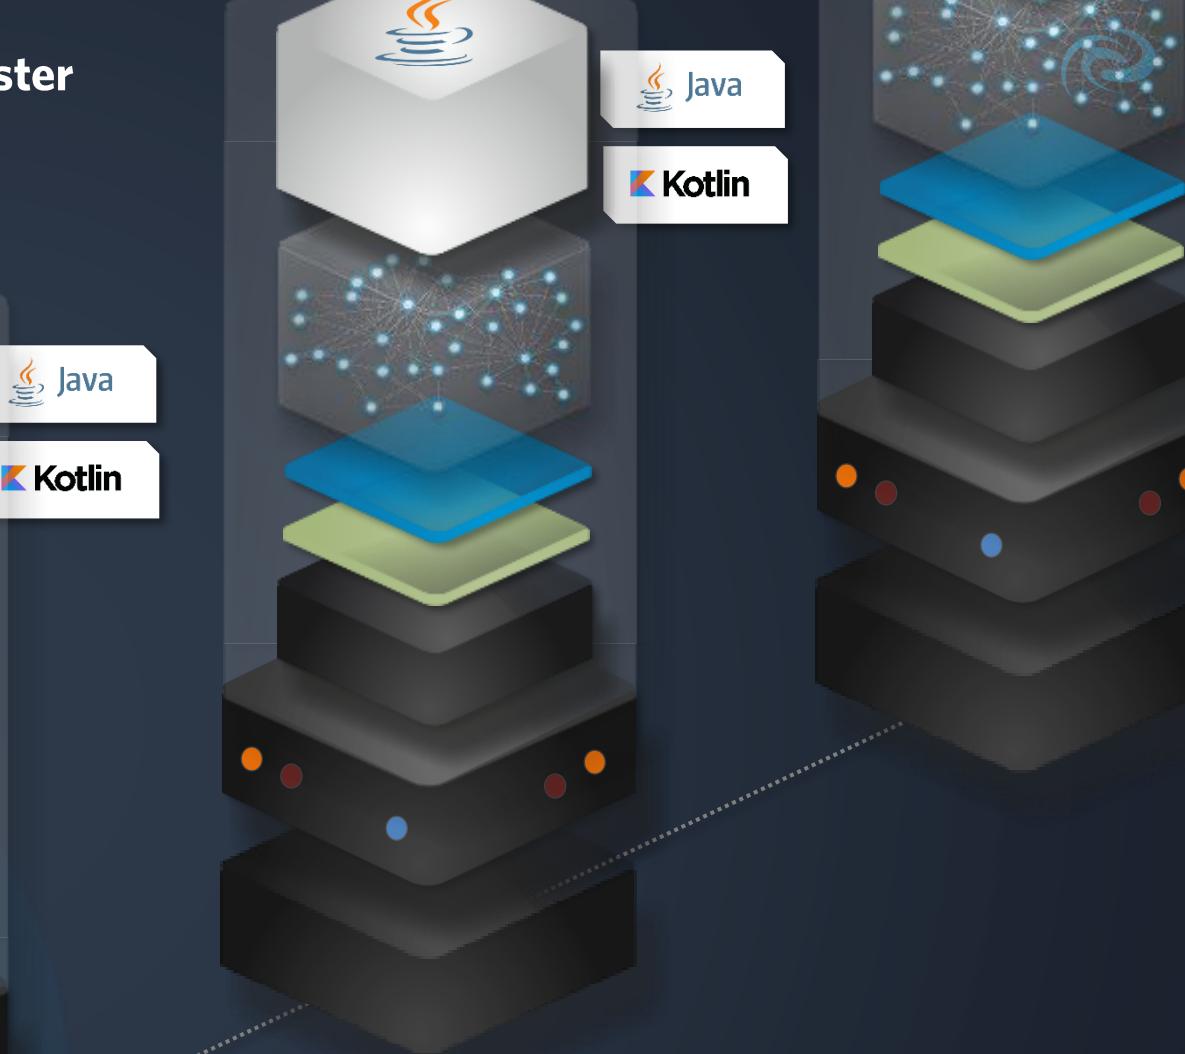
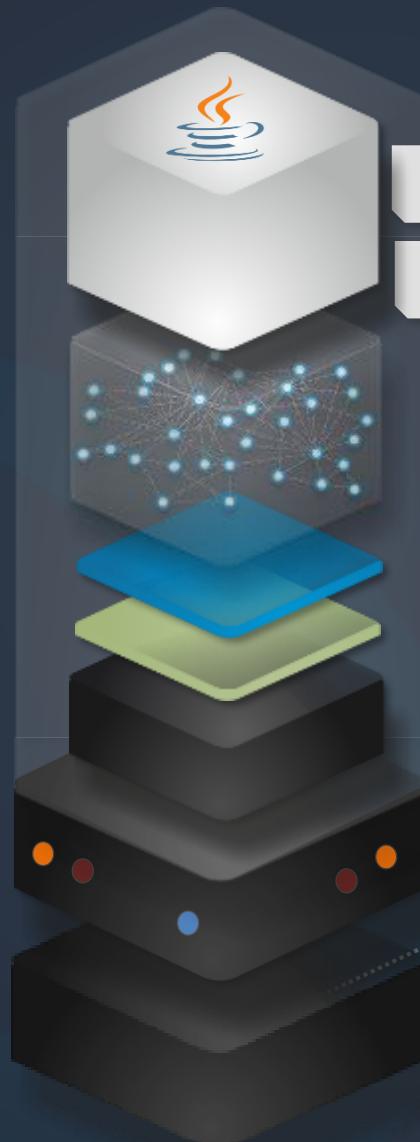
In-Memory Data (Java Object Graph)

**EclipseStore (Storage Engine)
Any Micro Runtime**

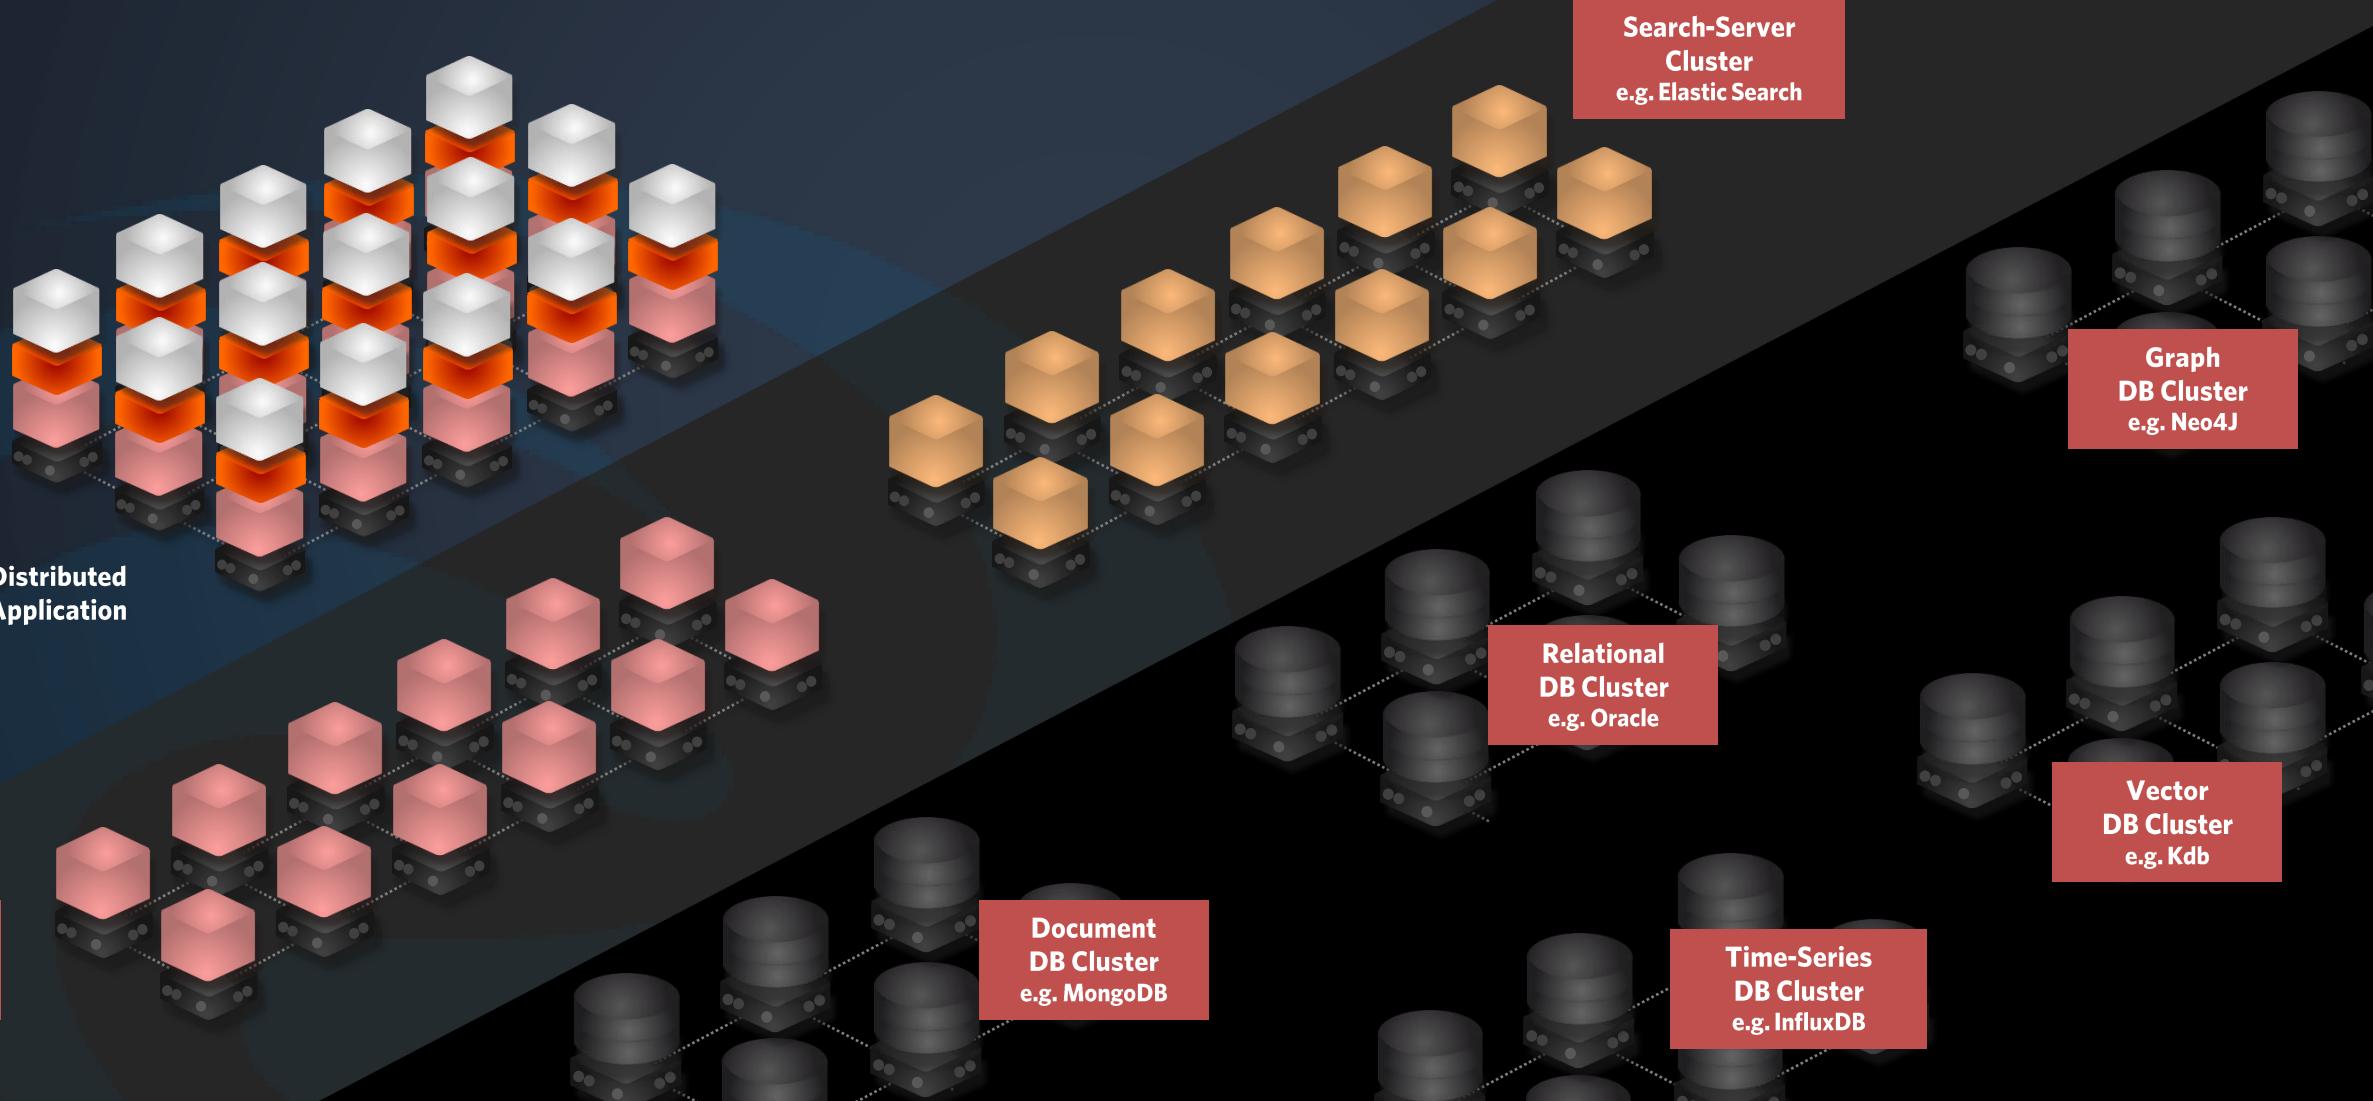
JVM

VM / Container / Pod

**Storage
Full Copy of All Data
(HDD/SSD/Persistent Volume,
BLOB Storage)**



From Today's Expensive Data Solutions Djungle to ...



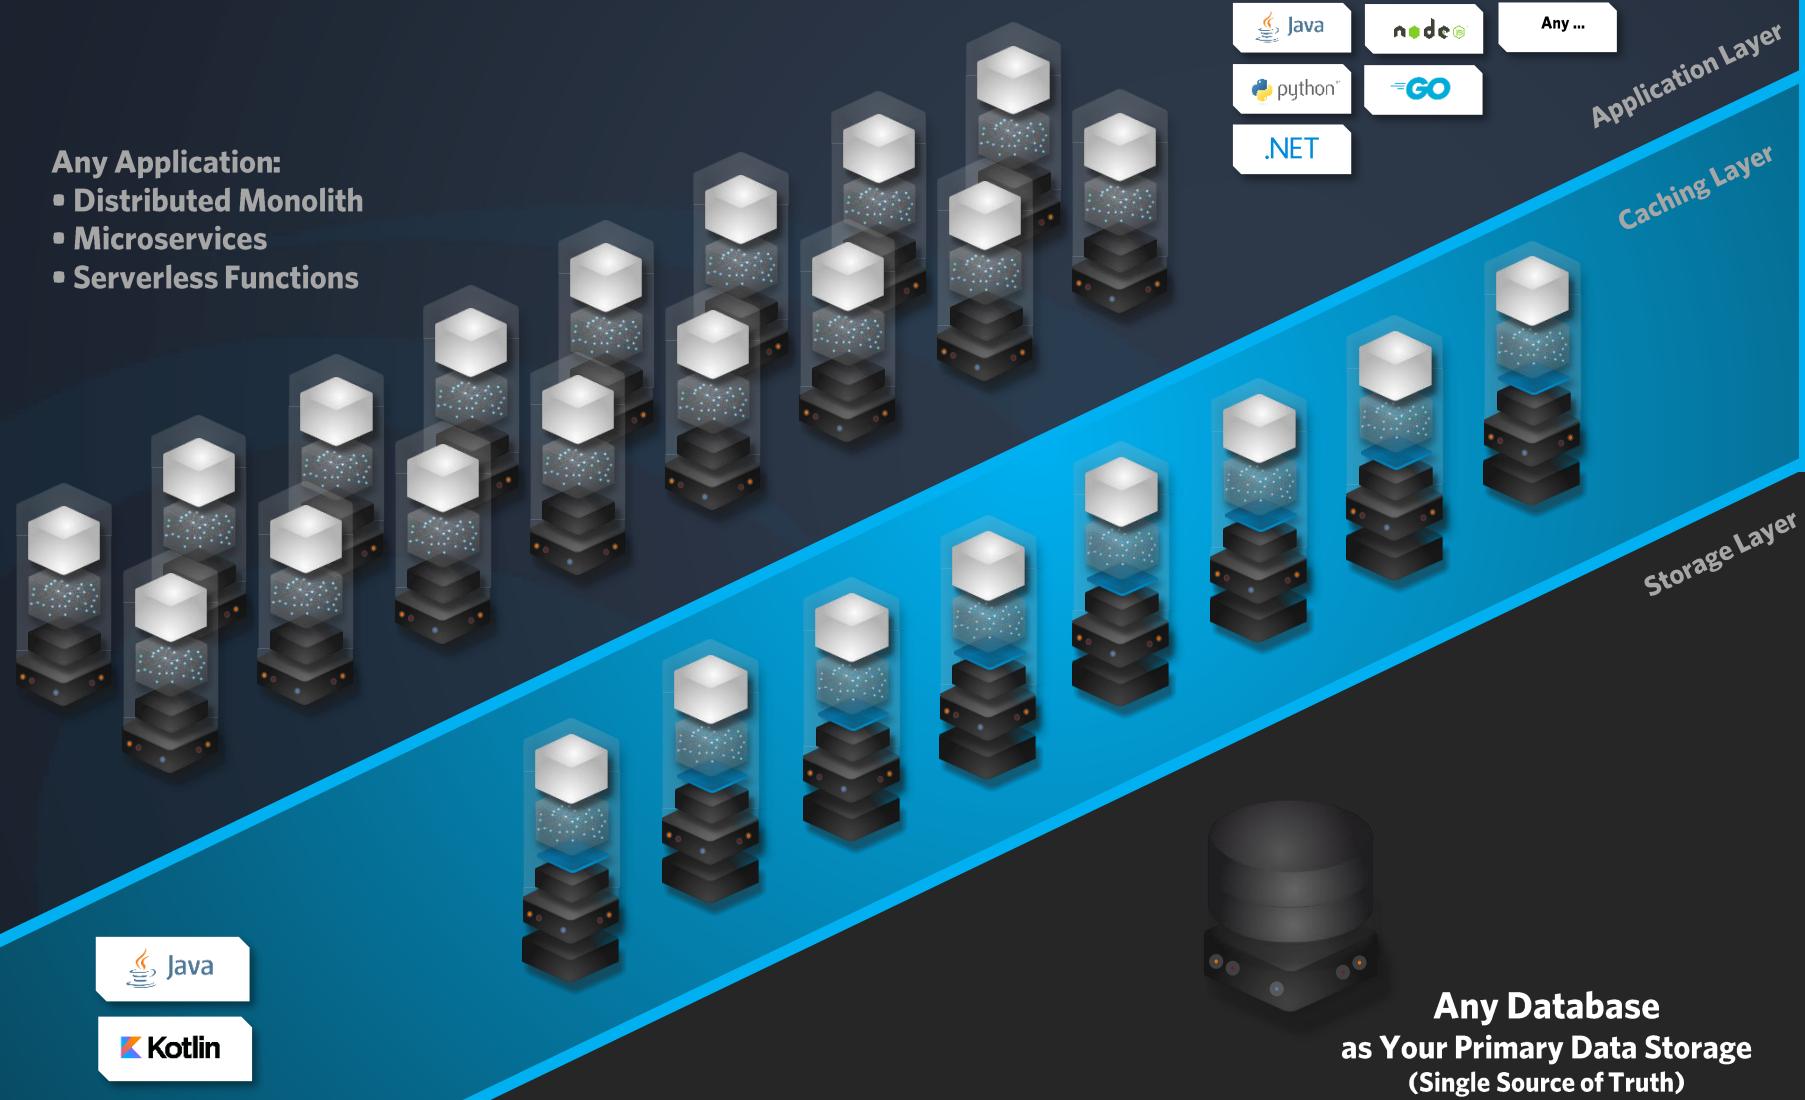


Eclipse Data Grid: Caching & In-Memory Searching Layer - and Beyond.

Speed-Up Any Database Application & Benefit From Significant Cost Savings: Compute. Storage. Development. Licenses.

Any Application:

- Distributed Monolith
- Microservices
- Serverless Functions



Eclipse Data Grid

- Caching
- Searching
- Indexing
- Native Java Object Model
- Any Logic in Java
- ACID Persistence

1000x

Faster Queries In-Memory



IBM Cloud

Eclipse Data Grid: As Primary Data Storage.

Speed-Up Any Database Application & Benefit From Dramatic Cost Savings: Compute. Storage. Development. Licenses.

- Any Application:
- Distributed Monolith
 - Microservices
 - Serverless Functions



Note: This architecture is most powerful and efficient, but in practice it is mostly refused by traditional database teams which is the biggest show-stopper for this solution.

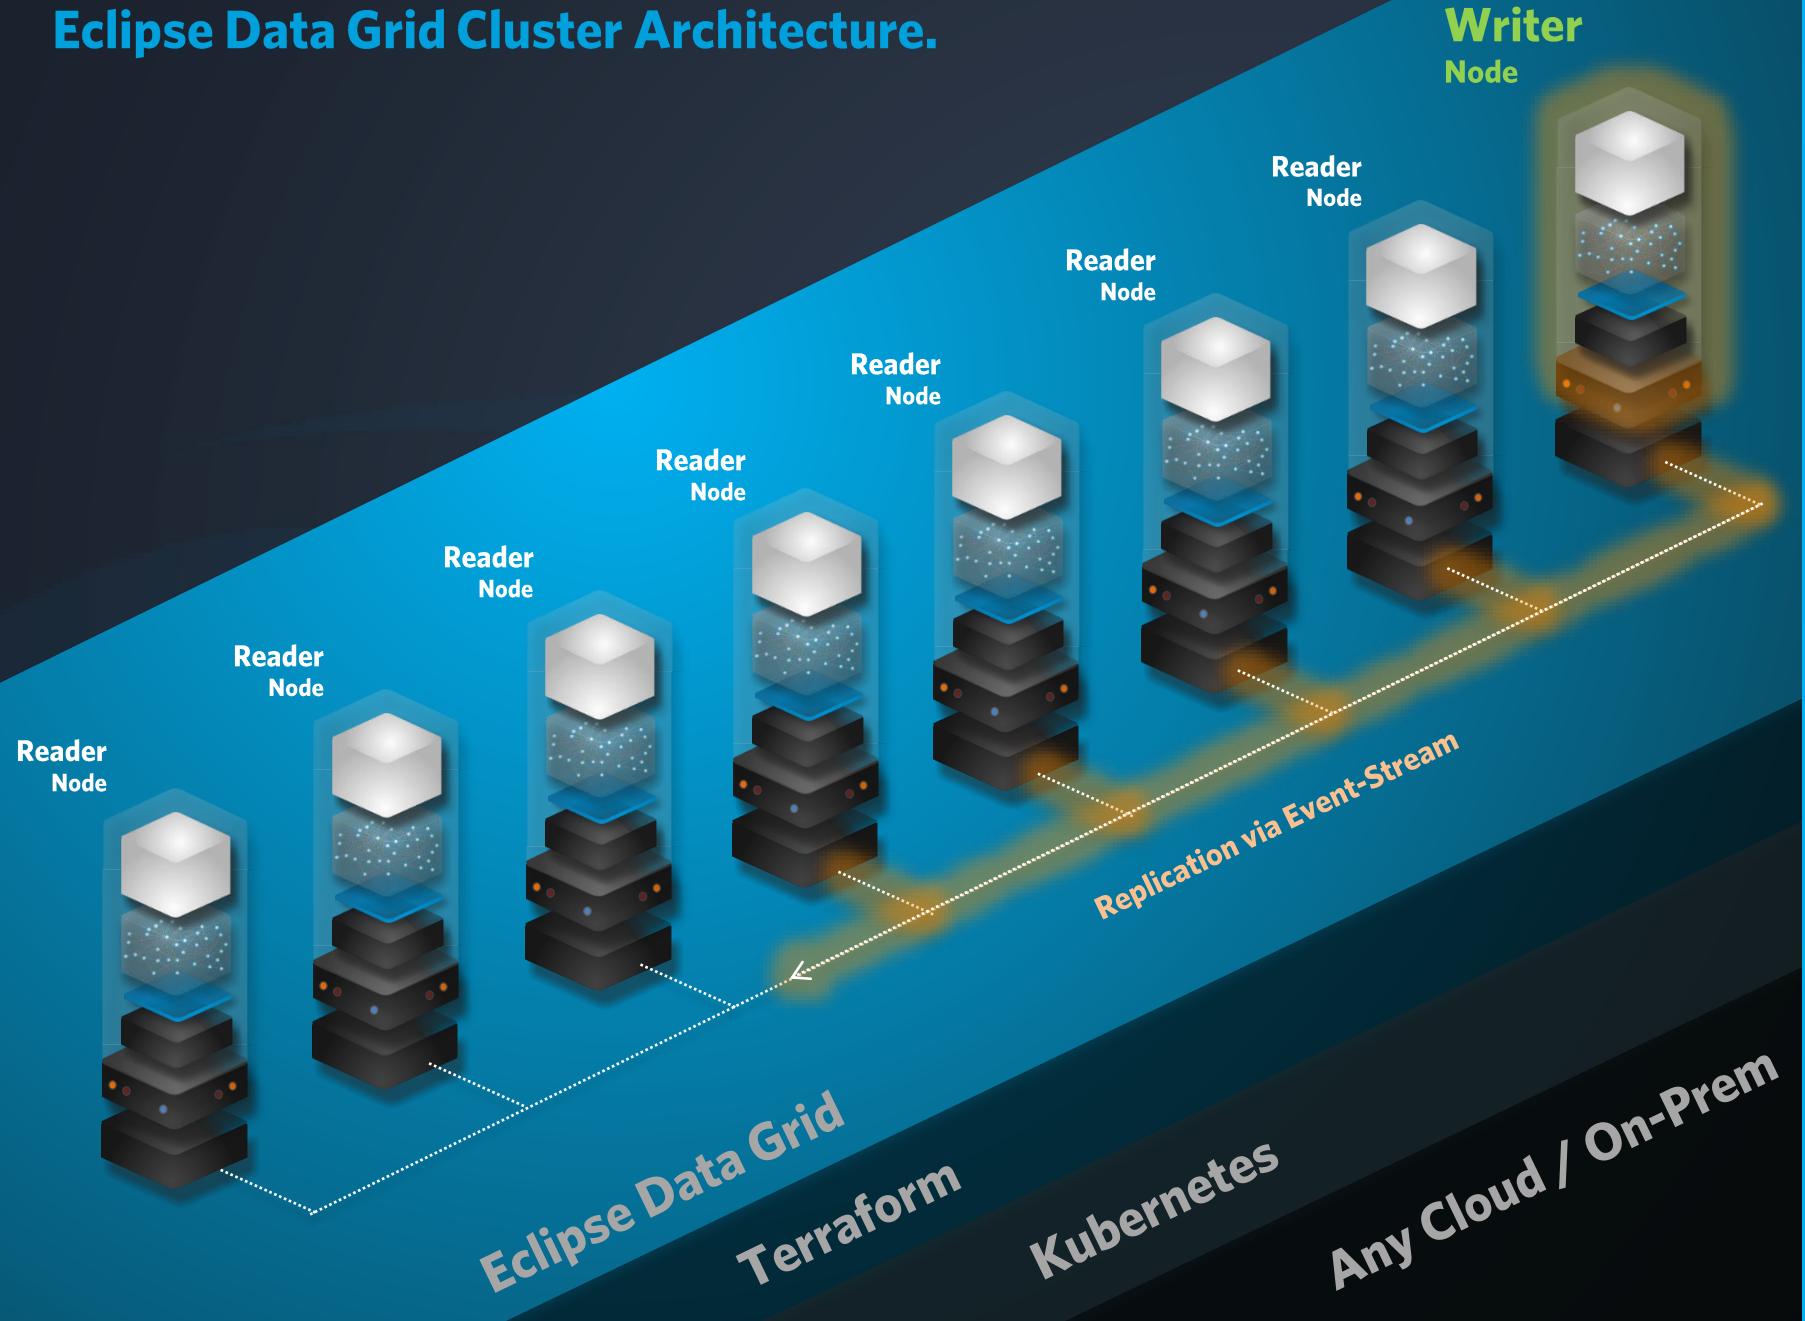
1000x
Faster Queries In-Memory

93%
Database Cost Savings

99%
Energy & CO₂ Emission Savings

Eclipse Data Grid Inside.

Eclipse Data Grid Cluster Architecture.



Eclipse Data Grid

Single Writer

Replication

Auto Fail-Over

Data Redundancy

High Availability

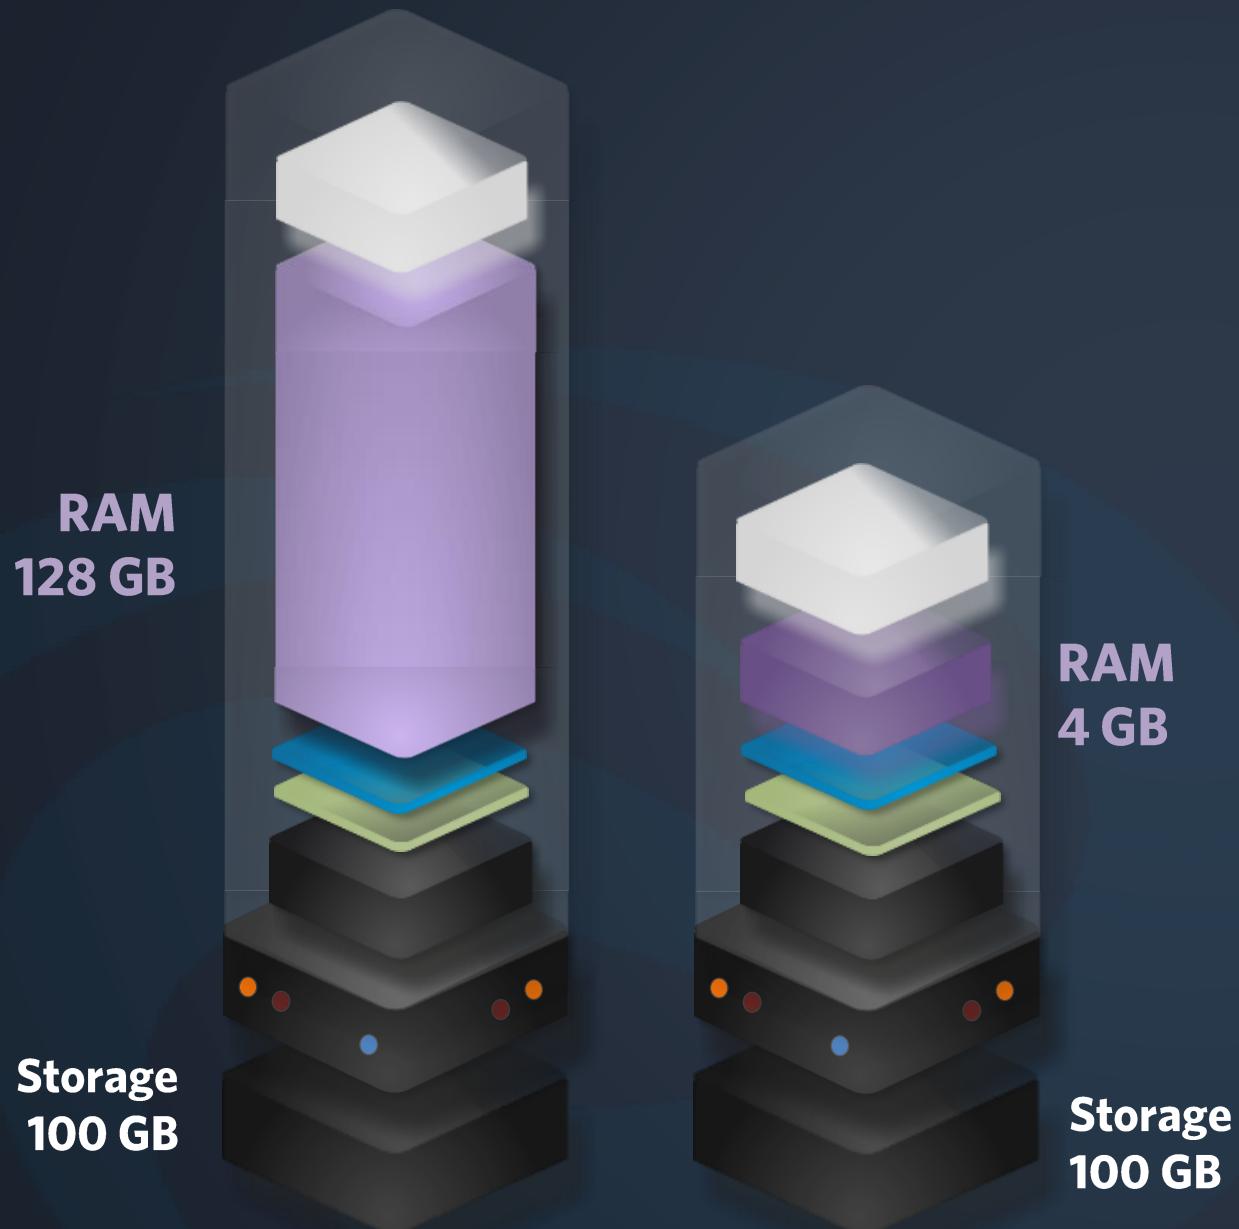
Elastic Scalable

Self Healing

Multi Region Backup

Consistency Levels:

- Local Node: Strong Consistency
- Writer Node: Strong Consistency
- Cluster:
 - Standard: Eventual Consistency
 - Customizable: Strong Consistency



**Lazy-Loading Retrieving Data from SSD Drives is Enormously Fast:
Mostly Millisecond Latencies**

**Full In-Memory Computing is the
Fastest Approach.**

	Small RAM	Big RAM
IO Ops	More	Less
Performance	Lower	Higher
Start-Up	Faster	Slower
Infrastructure Costs	Lower	Higher



The OpenJ9 Advantage

- Superior RAM Footprint & Efficiency: Up to 65% Smaller
- Advanced Garbage Collection for In-Memory Performance
- Faster Ramp-Up to Peak Throughput
- JIT Optimization Levels



The OpenJ9 logo is displayed in a large, white, sans-serif font. The letters 'Open' are positioned above the letter 'J9', which is enclosed within a circular outline. The background of the slide features a dark, abstract digital circuit board or network diagram with glowing blue lines and nodes, creating a futuristic and technological atmosphere.

Eclipse Data Grid – Enterprise-Grade Open Source.



Open Source

Eclipse Data Grid

In-Memory Data Processing Platform

EclipseStore

Java Native Persistence

Eclipse Serializer

High-Secure Java Serialization

Eclipse OpenJ9 / Any JVM

Java Virtual Machine

Recommended, but not Mandatory!



Enterprise

SaaS

Multi - Cloud

- ✓ Managed
- ✓ Built-In Security
- ✓ Serverless
- ✓ Management UI
- ✓ Multi-Cloud

On-Prem

Requires only Kubernetes

- ✓ Tested Builds
- ✓ Full Ownership
- ✓ Data Sovereignty
- ✓ Customizing

✓ Enterprise Support

✓ Consulting, Implementation, Partner Support & Mentoring





Move Complex Data Processing & Manipulation to a Java In-Memory Layer



Why Eclipse Data Grid?

High Performance:

Microsecond Response- & Query-Time

Data Model:

Native Java Object Model

Multi-Data-Model:

Supporting Any Java Type, any Document, any Structure, any Complexity

Pure Java:

Write Query & Any Business Logic in Pure Java

Persistence:

ACID-Compliant Persistence. Can Be Used as a Primary Data Storage

Distributed:

Distributed, Replication, Data Redundancy

Security:

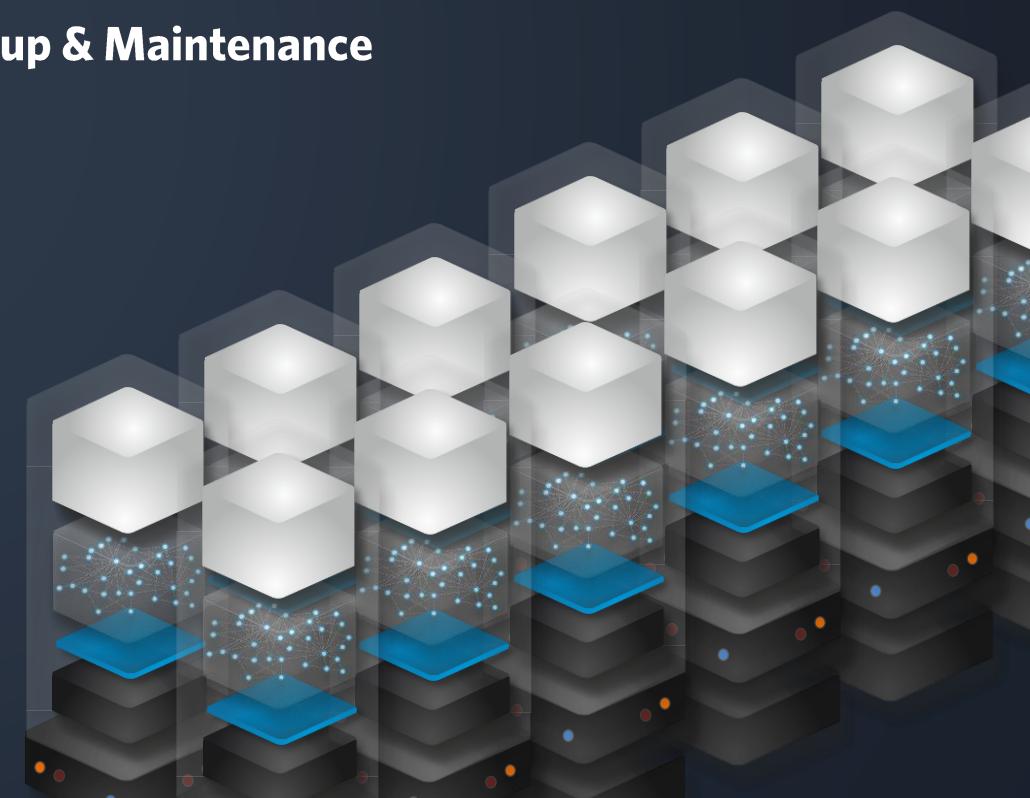
Built-In Security

Simplicity:

Convenient & Easy Setup & Maintenance

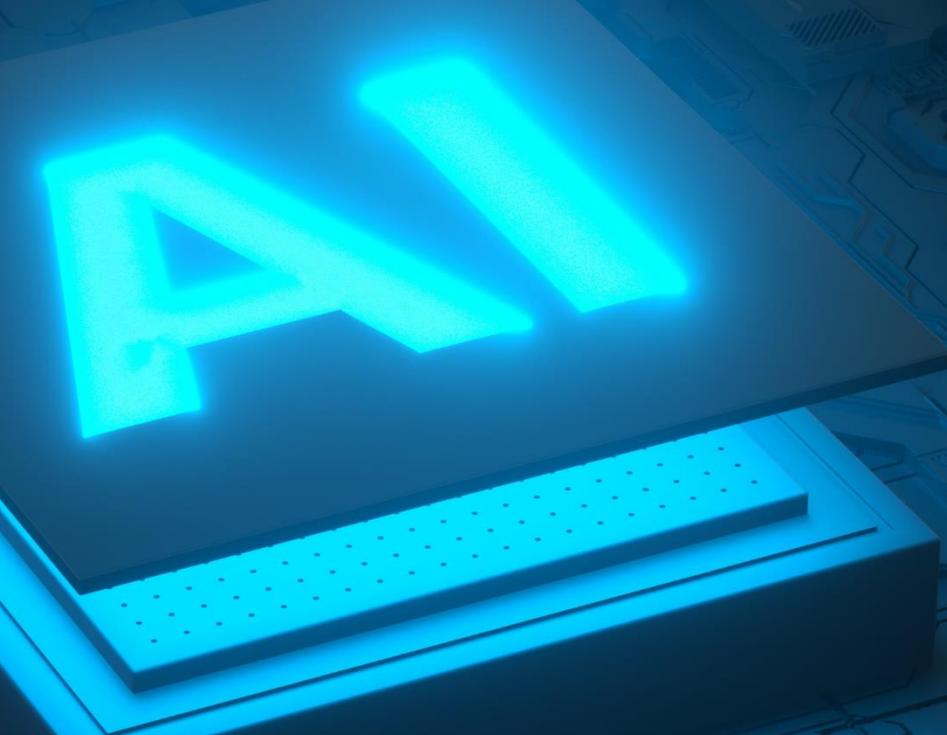
Available Platforms:

SaaS & On-Prem



What's Next:

Java Vector Search to Build Ultra-Fast Embedded In-Memory Vector Database Applications with Pure Java.





What's Next:

Vector Search

NEW !

Cyrock.AI: Java-Native Embedded In-Memory Vector Database.

VectraLink

NEW !

Synchronizing Traditional Databases with Your Java-Native Vector Database.

Eclipse Data Grid Features

Indexing

Extreme Fast Off-Heap Bitmap Index to Minimize IO Operations.

GigaMap

Powerful Map Enabling Fully Automated Lazy-Loading.

Locking API

Enabling Locking and Simplifying Concurrency Handling on the Java Object Graph Level.

Lucene Integration

Extends Eclipse Data Grid With Full-Text Search Capabilities.

ACID Persistence

Persisting Any Java Object Graph Transaction-Safe In Various Storage Targets.

Highly-Secure Serialization

Persisting Any Java Object Graph Transaction-Safe In Various Storage Targets.

Object Graph Replication

Distributing Any Java Object Graph Through Multiple JVMs in a Cluster.

Cluster

Configuration Files For Creating an Eclipse Data Grid Cluster Environment Based on Kubernetes.



Cyrock.AI

Embedded Java In-Memory Vector Database

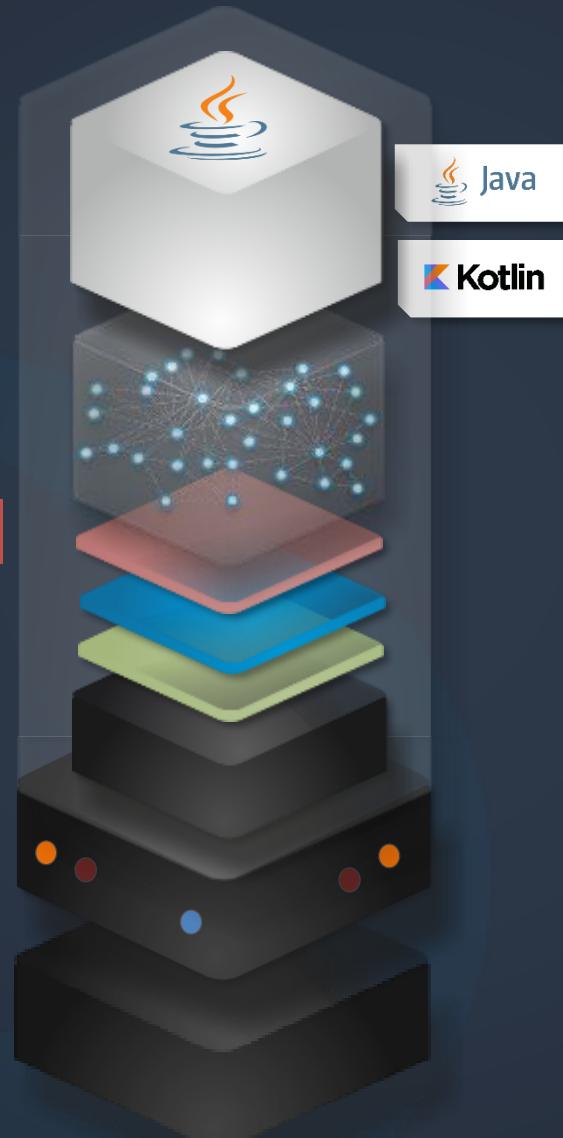
Queries in Java
Any Business Logic in Java

In-Memory Data (Java Object Graph)

Java Vector Search
EclipseStore (Storage Engine)
Micro Runtime

JVM
VM / Container / Pod

Storage
Full Copy of All Data
(HDD/SSD/Persistent Volume)



**The Classic Java Monolith +
EclipseStore:
Ultra-Fast Java In-Memory
Data Monster.**

Cyrock.AI VectraLink:

Business Data

Business Data

Business Data

Business Data

Business Data

Business Data



Chroma



VectraLink

Automated Synchronization of Vector Embeddings between Enterprise & Vector Databases



Vector DB

Vector DB

Vector DB

Vector DB

Vector DB

Vector DB



 CYROCK.AI

Eclipse Data Grid

Java In-Memory Data Platform

<https://www.microstream.one>

EclipseStore

Java Native Persistence

<https://www.eclipsestore.io>



CYROCK.AI

**Embedded Java In-Memory
Vector Database**

<https://github.com/cyrock-ai/early-access/>

Contact:

Markus Kett

Co-Founder & CEO

m.kett@cyrock.ai

Cell: +49 16090766677

LinkedIn: MarkusKett



CYROCK.AI

CYROCK.AI, Inc.

info@cyrock.ai

www.cyrock.ai